

18-640 Foundations of Computer Architecture

Lecture 3: “From Pipelined to Superscalar Processors”

John Paul Shen

September 2, 2014

➤ Required Reading Assignment:

- Chapters 3 and 4 of Shen and Lipasti (SnL).

➤ Recommended Reading Assignment:

- “*The Interaction of Architecture and Operating System Design*,” by Anderson, Levy, Bershad, and Lazowska, ASPLOS 1991.

9/2/2014 (© J.P. Shen)

18-640 Lecture 3



Carnegie Mellon University 1

18-640 Foundations of Computer Architecture

Lecture 3: “From Pipelined to Superscalar Processors”

- A. Limitations of Scalar Pipelined Processors
 - IBM RISC Experience
- B. Modern Superscalar Processor Organization
 - Parallel, Diversified, and Dynamic Pipelines
- C. Limits on Instruction Level Parallelism (ILP)
 - Motivation for Superscalar Processors
- D. Computer Architecture & Operating Systems

9/2/2014 (© J.P. Shen)

18-640 Lecture 3

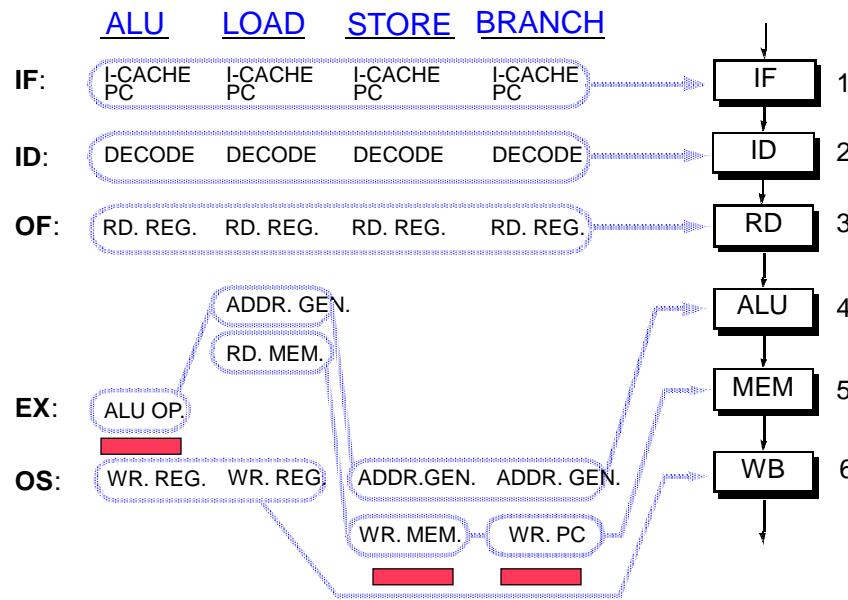


Carnegie Mellon University 2

Review from Lecture 2

Unifying Instruction Types in a Pipelined Processor

The 6-stage TYPICAL pipeline:



9/2/2014 (© J.P. Shen)

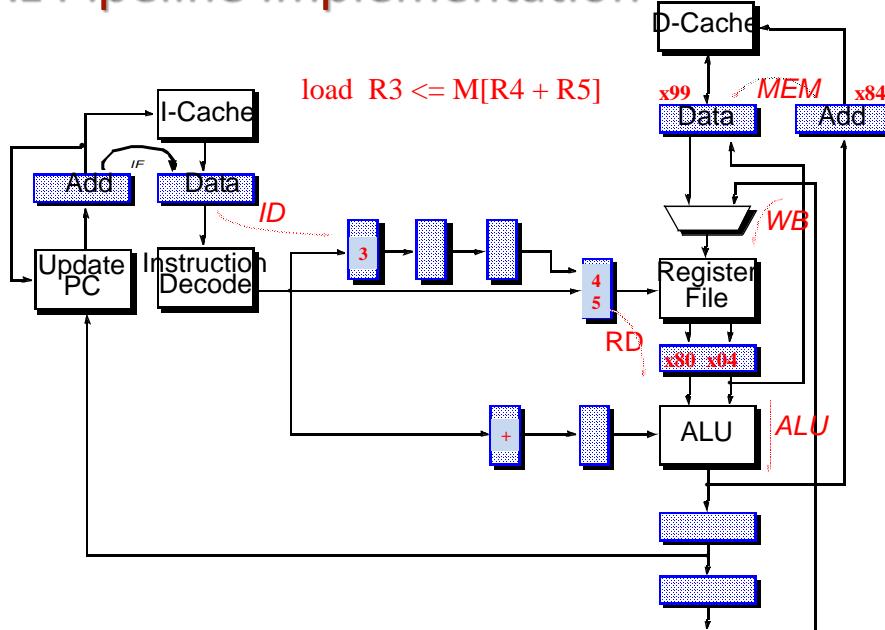
18-640 Lecture 3

Carnegie Mellon University 3

Review from Lecture 2

TYPICAL Pipeline Implementation

The 6-stage TYPICAL pipeline:



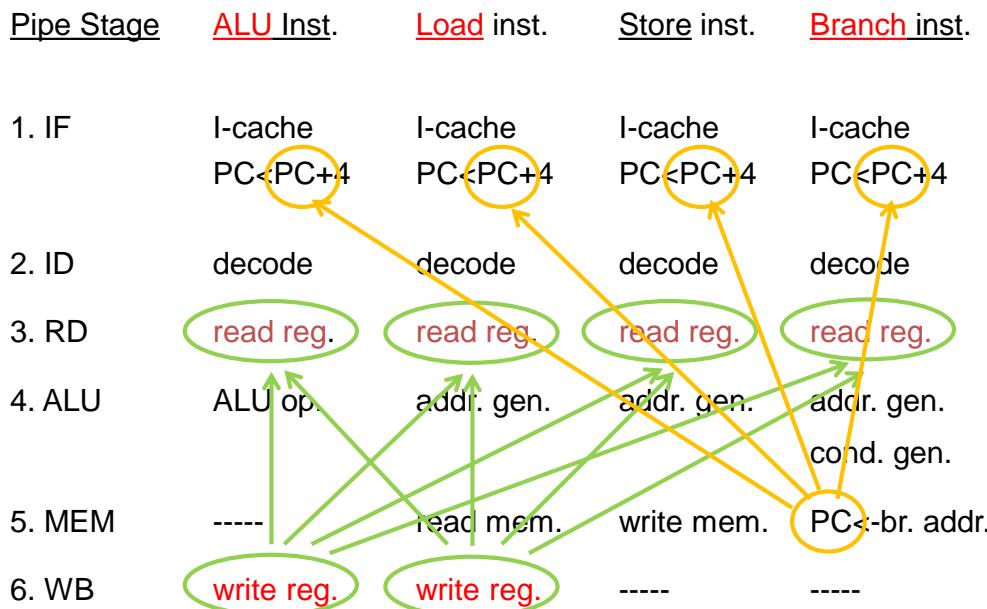
9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 4

Review from Lecture 2

Inter-Instruction Hazards



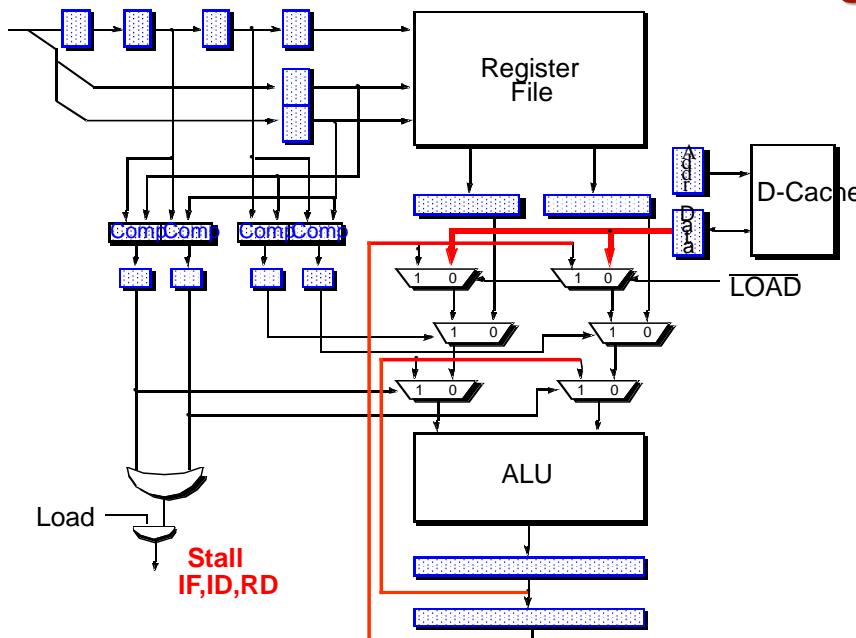
9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 5

Review from Lecture 2

Implementation of ALU and Load Forwarding Paths



9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 6

Review from Lecture 2

Penalties with Forwarding Paths

Leading Inst _i	ALU	Load	Branch
Trailing Inst _j	ALU, L/S, Br.	ALU, L/S, Br.	ALU, L/S, Br.
Hazard register	Int. Reg. (R _i)	Int. Reg. (R _i)	PC
Register WRITE stage (i)	WB (stage 6)	WB (stage 6)	MEM (stage 5)
Register READ stage (j)	RD (stage 3)	RD (stage 3)	IF (stage 1)
Forward from outputs of:	ALU, MEM, WB	MEM, WB	MEM
Forward to input of:	ALU	ALU	IF
RAW distance or penalty:	0 cycles	1 cycles	4 cycles

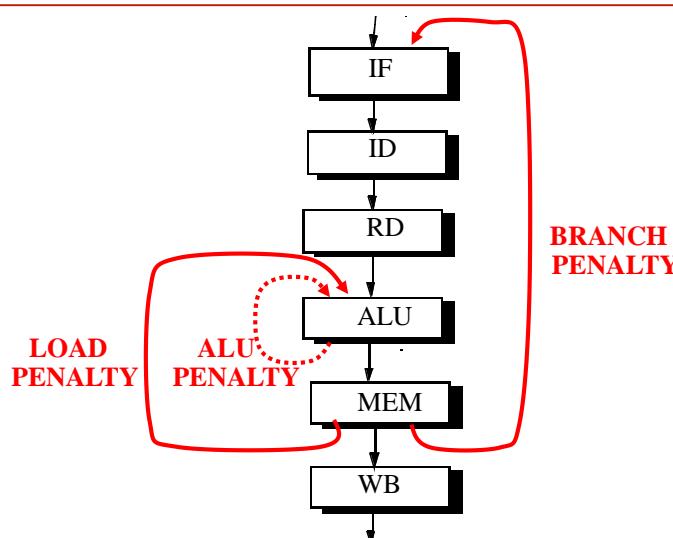
9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 7

Review from Lecture 2

3 Major Penalty Loops of (Scalar) Pipelining



Performance Objective: Reduce CPI as close to 1 as possible.

9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 8

IBM RISC Experience [Agerwala and Cocke 1987]

- Internal IBM study: Limits of a scalar pipeline?
- Memory Bandwidth
 - Fetch 1 instr/cycle from I-cache
 - 40% of instructions are load/store (D-cache)
- Code characteristics (dynamic)
 - Loads – 25%
 - Stores 15%
 - ALU/RR – 40%
 - Branches – 20%
 - 1/3 unconditional (always taken)
 - 1/3 conditional taken, 1/3 conditional not taken

A. Limitations of Scalar Pipelined Processors

- IBM RISC Experience
 - Control and data dependences add 15%
 - Best case CPI of 1.15, IPC of 0.87
 - Deeper pipelines (higher frequency) magnify dependence penalties
- This analysis assumes 100% cache hit rates
 - Hit rates approach 100% for some programs
 - Many important programs have much worse hit rates

Address Limitations of Scalar Pipelined Processors

- Upper Bound on Scalar Pipeline Throughput

Limited by IPC = 1

→ Parallel Pipelines

- Inefficient Unification Into Single Pipeline

Long latency for each instruction

Hazards and associated stalls

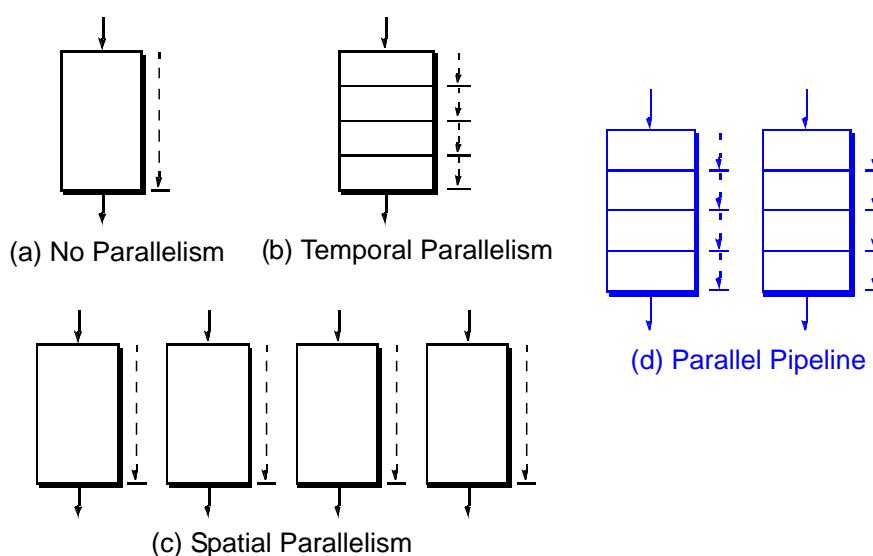
→ Diversified Pipelines

- Performance Lost Due to In-order Pipeline

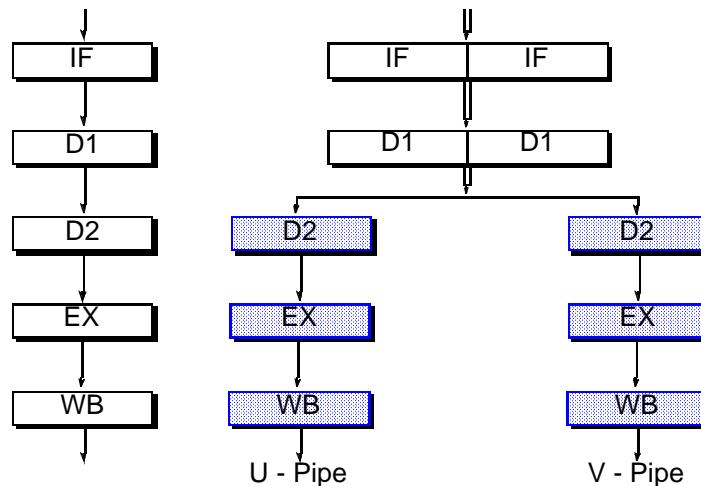
Unnecessary stalls

→ Dynamic Pipelines

Parallel Pipelines



Intel Pentium Parallel Pipeline

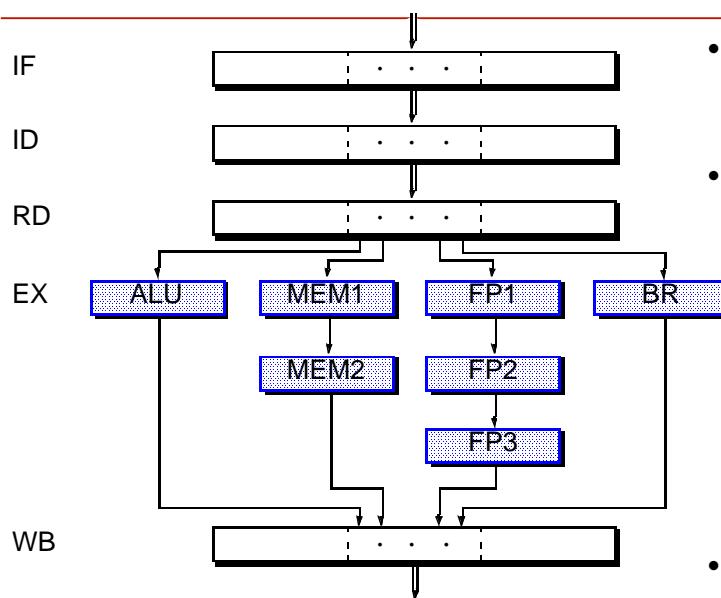


9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 13

Diversified Pipelines



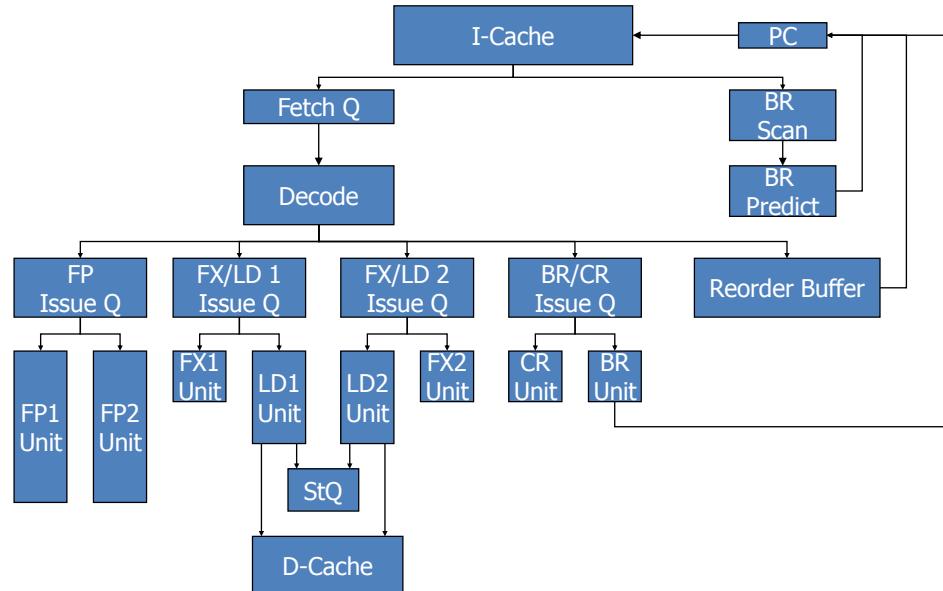
- Separate execution pipelines
 - Integer simple, memory, FP, ...
- Advantages:
 - Reduce instruction latency
 - Each instruction goes to WB asap
 - Eliminate need for forwarding paths
 - Eliminate some unnecessary stalls
 - E.g. slow FP instruction does not block independent integer instructions
- Disadvantages ??

9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 14

Power4 Diversified Pipelines

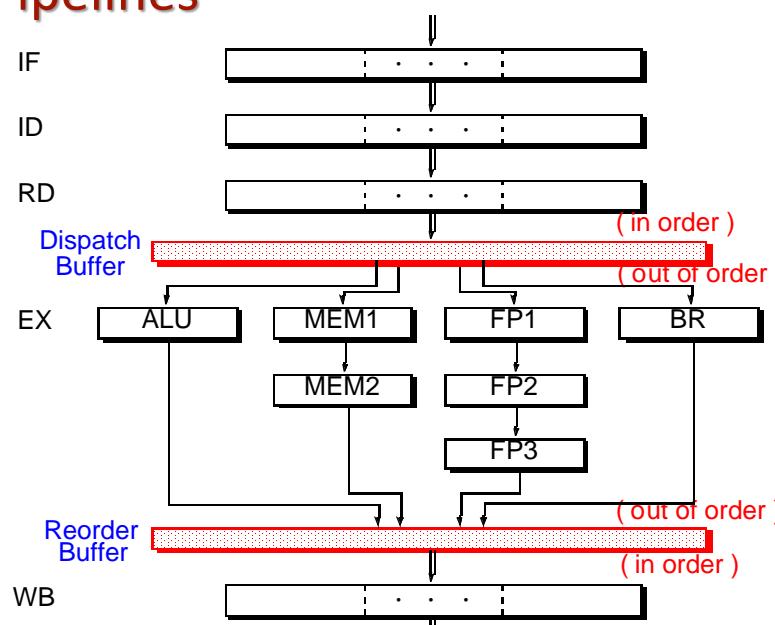


9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 15

Dynamic Pipelines

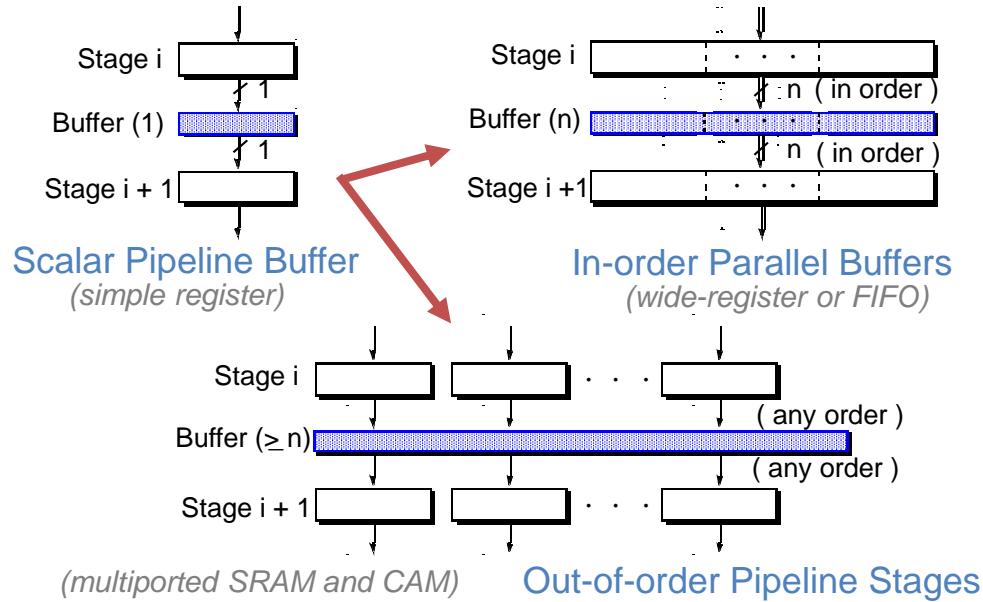


9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 16

Designs of Inter-stage Buffers

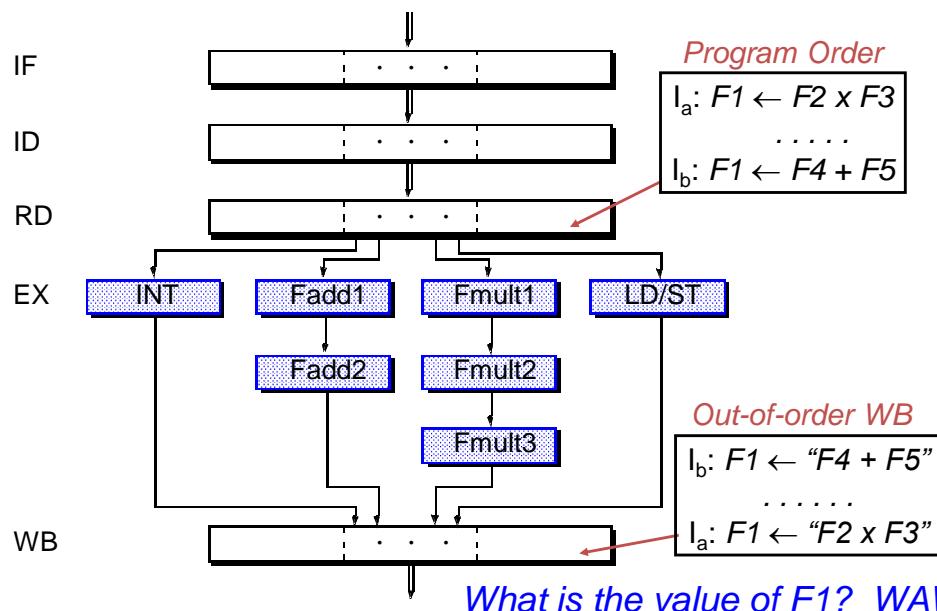


9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 17

The Challenges of Out-of-Order

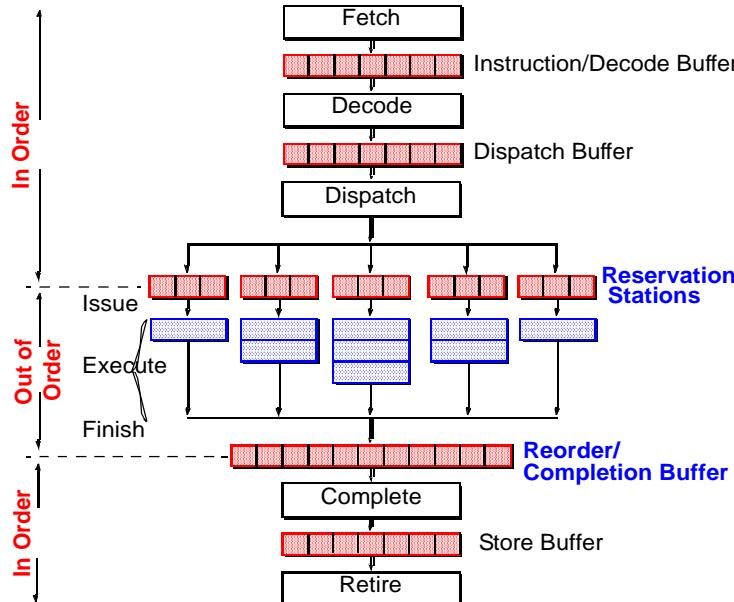


9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 18

B. Modern Superscalar Processor Organization

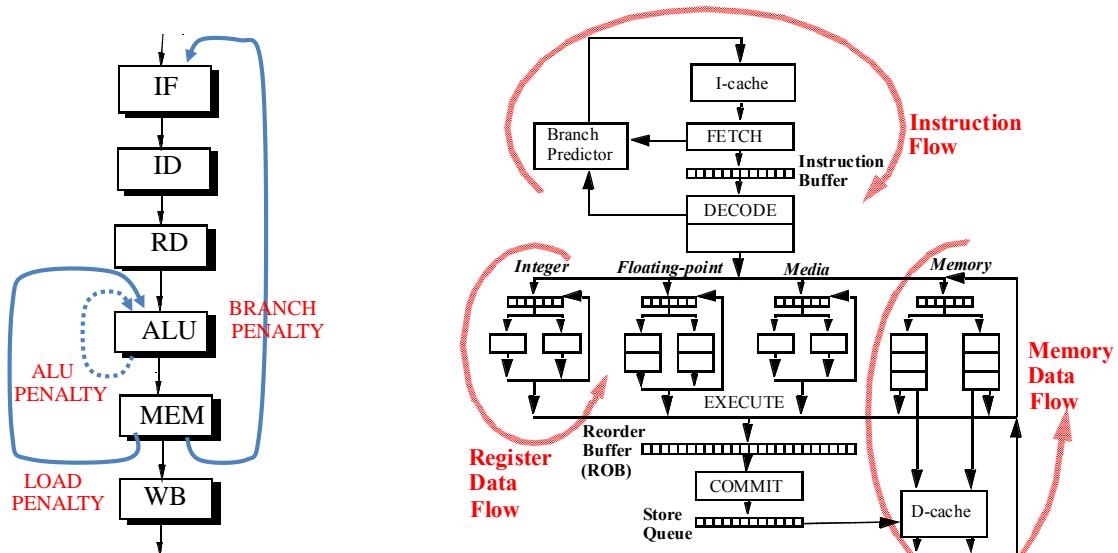


9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 19

Impediments to Superscalar Performance

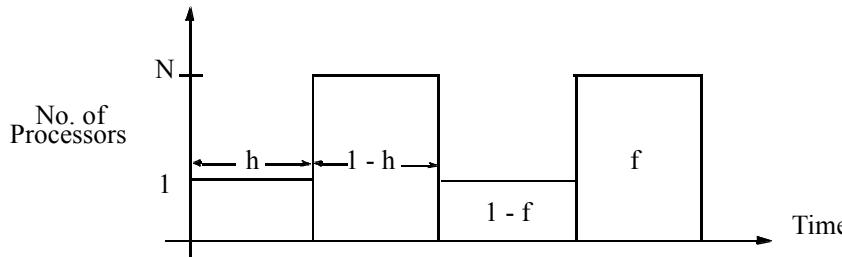


9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 20

Amdahl's Law and Instruction Level Parallelism



- h = fraction of time in serial code
- f = fraction that is vectorizable or parallelizable
- N = max speedup for f
- Overall speedup $\rightarrow \rightarrow$

$$\text{Speedup} = \frac{1}{(1-f) + \frac{f}{N}}$$

9/2/2014 (© J.P. Shen)

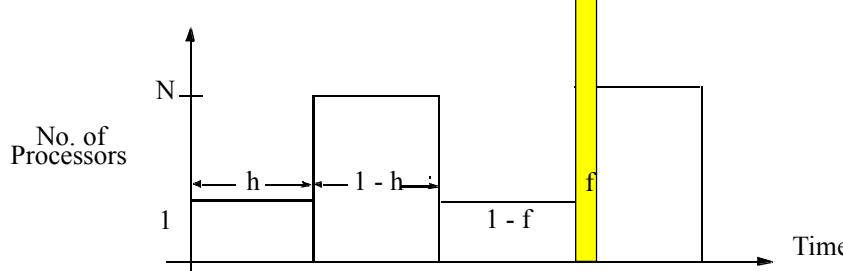
18-640 Lecture 3

Carnegie Mellon University 21

Revisit Amdahl's Law

- Sequential bottleneck
- Even if N is infinite
 - Performance limited by non-vectorizable portion ($1-f$)

$$\lim_{N \rightarrow \infty} \frac{1}{(1-f) + \frac{f}{N}} = \frac{1}{1-f}$$

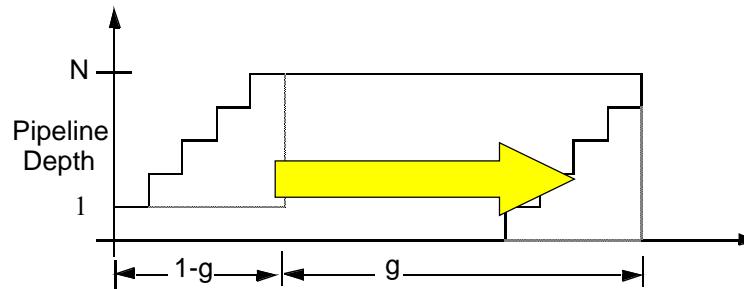


9/2/2014 (© J.P. Shen)

18-640 Lecture 3

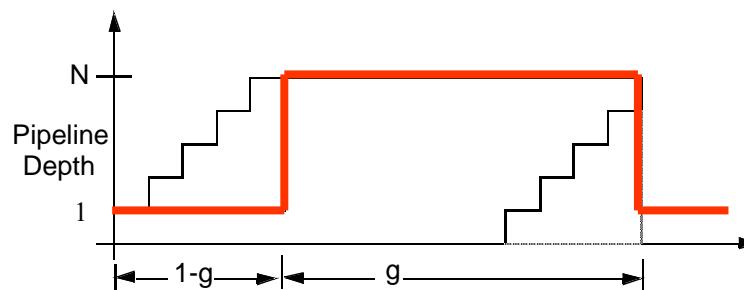
Carnegie Mellon University 22

Pipelined Processor Performance Model



- g = fraction of time pipeline is filled
- $1-g$ = fraction of time pipeline is not filled (stalled)

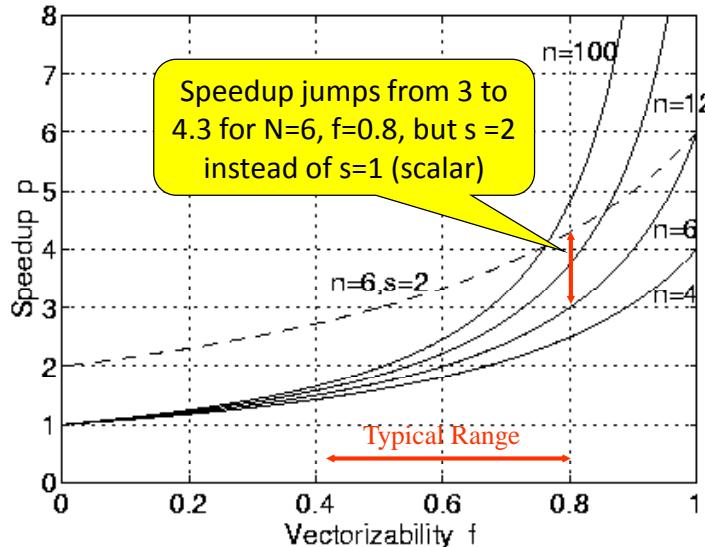
Pipelined Processor Performance Model



- “Tyranny of Amdahl’s Law”
 - When g is even slightly below 100%, a big performance hit will result
 - Stalled cycles in the pipeline are the key adversary and must be minimized as much as possible
 - Can we somehow fill the pipeline bubbles (stalled cycles)?

Motivation for Superscalar Design

[Tilak Agerwala and John Cocke, 1987]



9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 25

Superscalar Proposal

- Moderate the tyranny of Amdahl's Law
 - Ease the sequential bottleneck
 - More generally applicable
 - Robust (less sensitive to f)
 - Revised Amdahl's Law:

$$\text{Speedup} = \frac{1}{(1-f)} + \frac{f}{N}$$

9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 26

The Ideas Behind Modern Superscalar Processors

- Superscalar or wide instruction issue
Ideal IPC = n (CPI = 1/n)
- Diversified pipelines
Different instructions go through different pipe stages
Instructions go through needed stages only
- Out-of-order or data-flow execution
Stall only on RAW hazards and structural hazards
- Speculation
Overcome (some) RAW hazards through prediction

And it all relies on: **Instruction Level Parallelism (ILP)**

Independent instructions within sequential programs

C. Limits on Instruction Level Parallelism (ILP)

Weiss and Smith [1984]	1.58
Sohi and Vajapeyam [1987]	1.81
Tjaden and Flynn [1970]	1.86 (Flynn's bottleneck)
Tjaden and Flynn [1973]	1.96
Uht [1986]	2.00
Smith et al. [1989]	2.00
Jouppi and Wall [1988]	2.40
Johnson [1991]	2.50
Acosta et al. [1986]	2.79
Wedig [1982]	3.00
Butler et al. [1991]	5.8
Melvin and Patt [1991]	6
Wall [1991]	7 (Jouppi disagreed)
Kuck et al. [1972]	8
Riseman and Foster [1972]	51 (no control dependences)
Nicolau and Fisher [1984]	90 (Fisher's optimism)

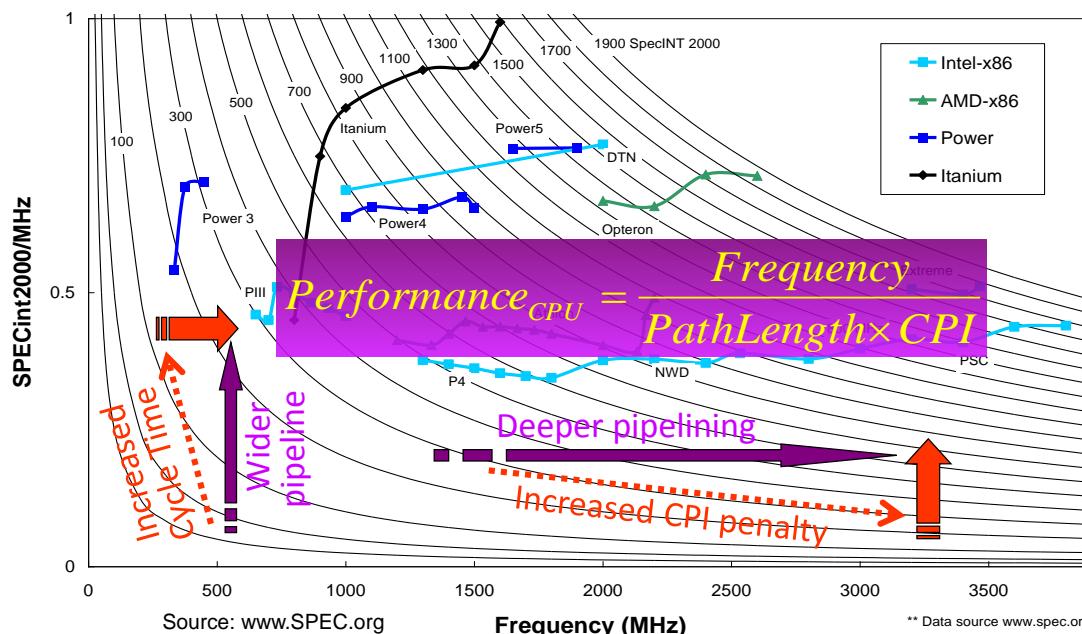
Processor Performance

$$\begin{aligned} \frac{1}{\text{Processor Performance}} &= \frac{\text{Time}}{\text{Program}} \\ &= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Time}}{\text{Cycle}} \\ &\quad (\text{path length}) \qquad \qquad (\text{CPI}) \qquad \qquad (\text{cycle time}) \end{aligned}$$

- In the 1980's (decade of **pipelining**):
 - ❖ CPI: 5.0 → 1.15
- In the 1990's (decade of **superscalar**):
 - ❖ CPI: 1.15 → 0.5 (best case)
- In the 2000's (decade of **multicore**):
 - ❖ Core CPI unchanged; chip CPI scales with #cores

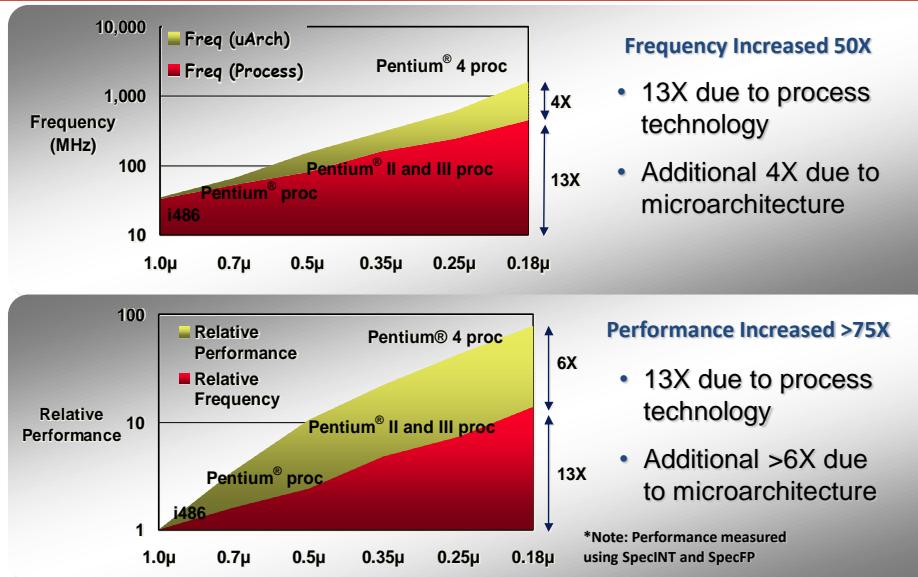
Frequency vs. Parallelism

[John DeVale & Bryan Black, 2005]



[Source: Intel Corporation]

Frequency and Performance Boost



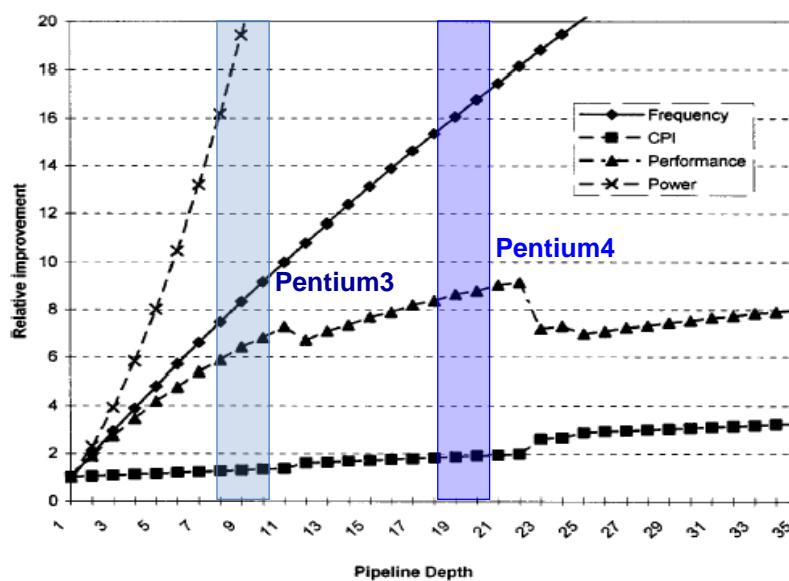
9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 31

[Ed Grochowski, Intel, 1997]

Putting it All Together: Limits to Deeper Pipelines



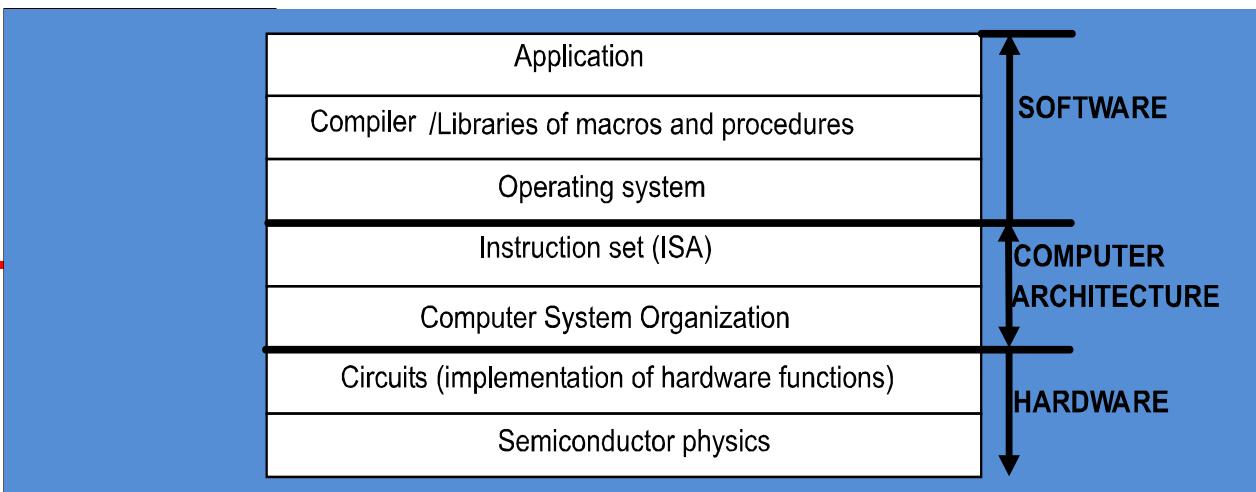
9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 32

[Hsien-Hsin Sean Lee, GIT 2007] [Xuxian Jiang, NCSU 2009]

D. Computer Architecture and Operating Systems



DSI = ISA = a contract between the program and the machine.

9/2/2014 (© J.P. Shen)

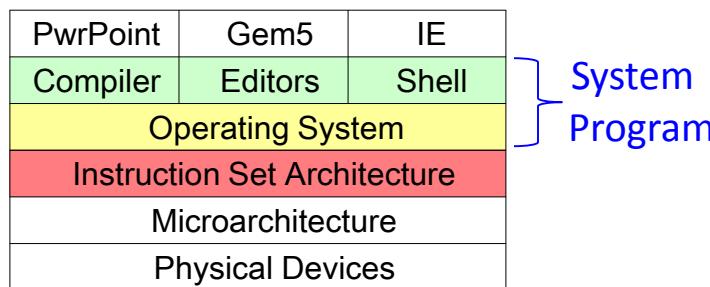
18-640 Lecture 3

Carnegie Mellon University 33

[Hsien-Hsin Sean Lee, 2007]

What is an Operating System?

- An intermediate program between a user of a computer and the computer hardware (to hide messy details)
- Goals:
 - Execute user programs and make solving user problems easier
 - Make the computer system convenient and efficient to use



9/2/2014 (© J.P. Shen)

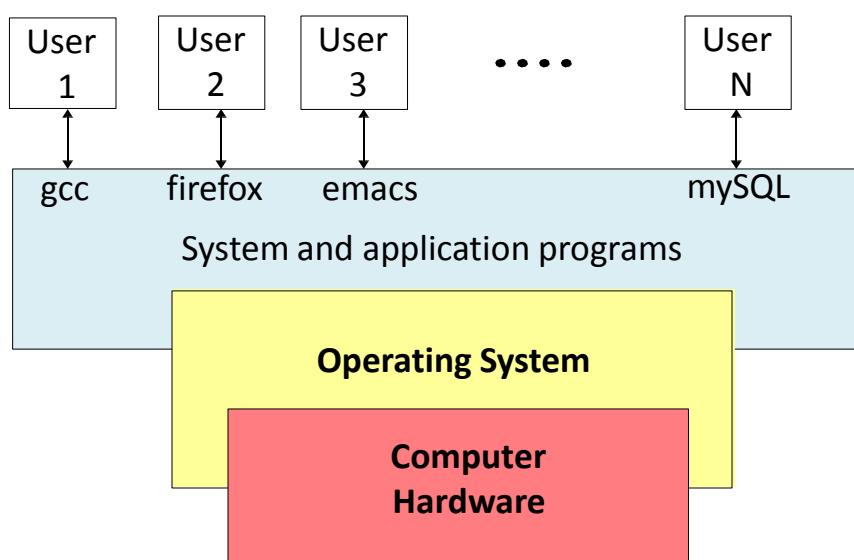
18-640 Lecture 3

Carnegie Mellon University 34

Computer System Components

- **Hardware**
 - Provides basic computing resources (CPU, memory, I/O)
- **Operating System**
 - Controls and coordinates the use of the hardware among various application programs for various users
- **Application Programs**
 - Define the ways in which the system resources are used to solve the computing problems of users (e.g. database systems, 3D games, business applications)
- **Users**
 - People, machines, other computers

Abstract View of Computer System Components



Time-Sharing Computing Systems

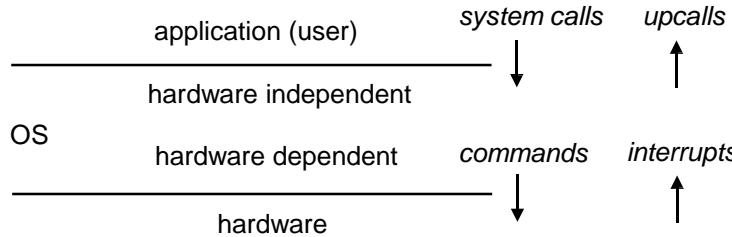
- CPU is multiplexed among several jobs that are kept in memory and on disk
(The CPU is allocated to a job only if the job is in memory)
- A job swapped in and out of memory to the disk
- On-line communication between the user and the system is provided
 - When the OS finishes the execution of one command, it seeks the next “control statement” from the user’s keyboard
- On-line system must be available for users to access data and code
- MIT MULTICS (MULTIplexed Information and Computing Services)
 - Ken Thompson went to Bell Labs and wrote one for a PDP-7
 - Brian Kernighan jokingly dubbed it UNICS (UNIplexed ..)
 - Later spelled to UNIX and moved to PDP-11/20
 - IEEE POSIX to standardize UNIX

Operating System Concepts

- Process Management
- Main Memory Management
- File Management
- I/O System Management
- Networking
- Protection System
- Command-Interpreter System

[Xuxian Jiang, NCSU 2009]

How Does an Operating System Work?

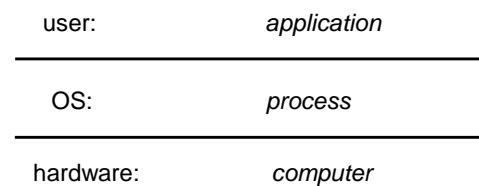


- Receives requests from the application: system calls
- Satisfies the requests: may issue commands to hardware
- Handles hardware interrupts: may upcall the application

- Abstraction
 - Process, memory, I/O, file, socket, ...
- Tradeoff
 - Separation between mechanisms and policies

Abstraction I: Process

A process is a system abstraction:
illusion of being the only job in the system



- Mechanism:
 - Creation, destruction, suspension, context switch, signalling, IPC, etc.
- Policy:
 - How to share system resources between multiple processes?

Abstraction II: Thread

A thread is a processor abstraction: illusion of having 1 processor per execution context

application:	<i>execution context</i>
OS:	<i>thread</i>
hardware:	<i>processor</i>

Process vs. Thread:
Process is the unit of resource ownership, while Thread is the unit of instruction execution.

➤ Mechanism:

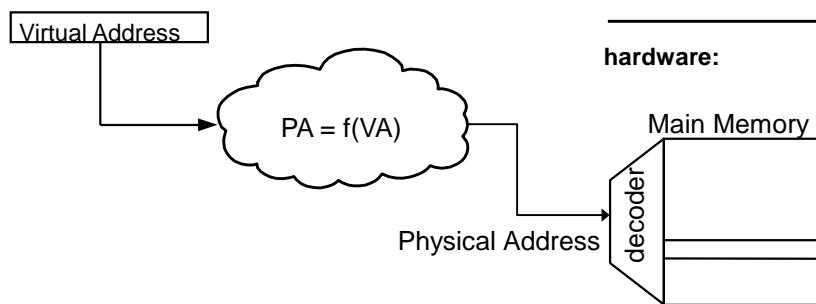
- Creation, destruction, suspension, context switch, signalling, synchronization, etc.

➤ Policy:

- How to share the CPU between threads from different processes?
- How to share the CPU between threads from the same process?

Abstraction III: Virtual Memory

Virtual memory is a memory abstraction:
illusion of large contiguous memory, often more
memory than physically available

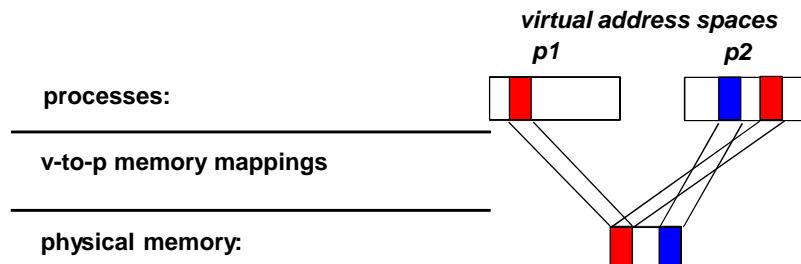


application:	<i>address space</i>
OS:	<i>virtual memory</i>
hardware:	<i>physical memory</i>

Virtual Memory Mechanism and Policy

➤ Mechanism:

- Virtual-to-physical memory mapping, page-fault, etc.



➤ Policy:

- How to multiplex a virtual memory that is larger than the physical memory onto what is available?
- How to share physical memory between multiple processes?

Abstraction IV: File System

A file system is a storage abstraction:
illusion of structured storage space

application/user: *copy file1 file2*

OS: *files, directories*

hardware: *disk*

➤ Mechanism:

- File creation, deletion, read, write, file-block- to-disk-block mapping, file buffer cache, etc.

➤ Policy:

- Sharing vs. protection?
- Which block to allocate for new data?
- File buffer cache management?

Abstraction V: Messaging

Message passing is a communication abstraction:
illusion of reliable (sometimes ordered) transport

application:	<i>sockets</i>
OS:	<i>TCP/IP protocols</i>
hardware:	<i>network interface</i>

- Mechanism:
 - Send, receive, buffering, retransmission, etc.
- Policy:
 - Congestion control and routing
 - Multiplexing multiple connections onto a single NIC

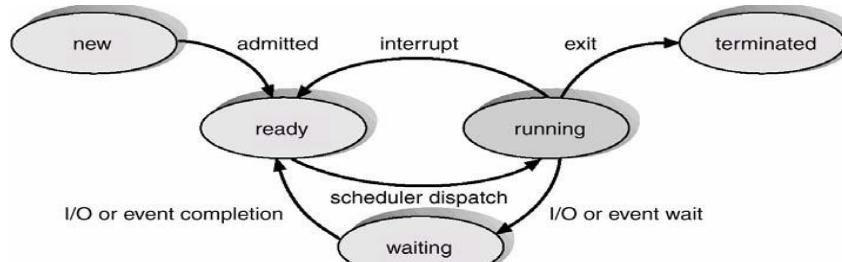
Process Management

- A process is a program in execution
- A process contains
 - Address space (e.g. read-only code, global data, heap, stack, etc)
 - PC, \$sp
 - Opened file handles
- A process needs certain resources, including CPU time, memory, files, and I/O devices
- The OS is responsible for the following activities for process management
 - Process creation and deletion
 - Process suspension and resumption
 - Provision of mechanisms for:
 - Process synchronization
 - Process communication

Process State

- As a process executes, it changes *state*:

- **New:** The process is being created
- **Ready:** The process is waiting to be assigned to a processor
- **Running:** Instructions are being executed
- **Waiting:** The process is waiting for some event (e.g. I/O) to occur
- **Terminated:** The process has finished execution



9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 47

Process Control Block (PCB)

Information associated with each process:

- Process state
- Program counter
- CPU registers (for context switch)
- CPU scheduling information (e.g. priority)
- Memory-management information (e.g. page table, segment table)
- Accounting information (PID, user time, constraint)
- I/O status information (list of I/O devices allocated, list of open files etc.)

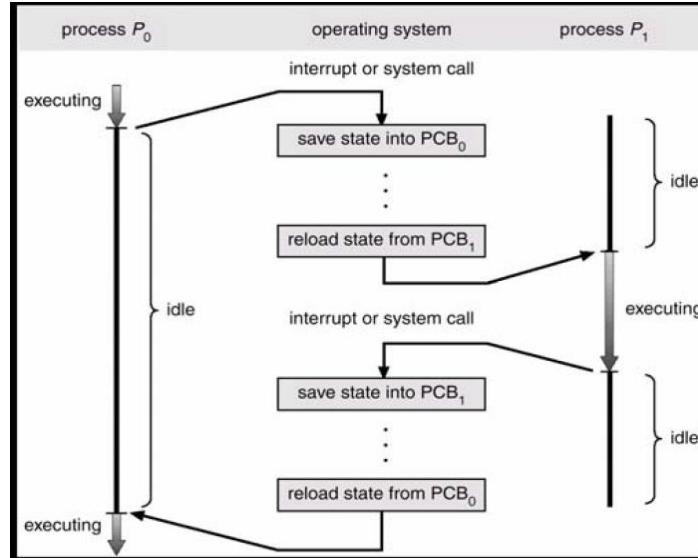
process state
process number
program counter
registers
memory limits
list of open files
...

9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 48

CPU Switches from Process to Process

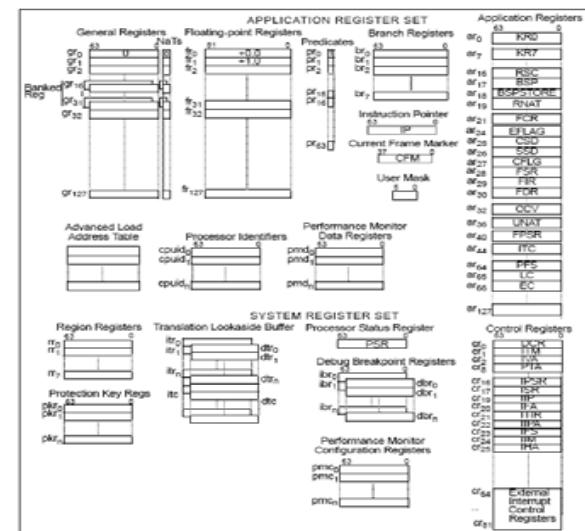
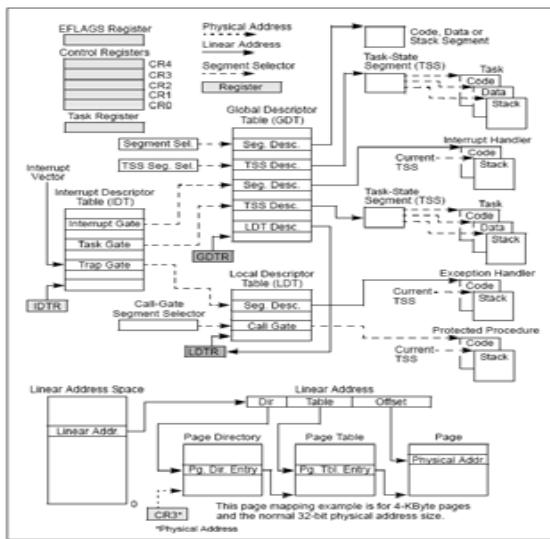


9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 49

Examples of System State: Intel Core Duo & Itanium



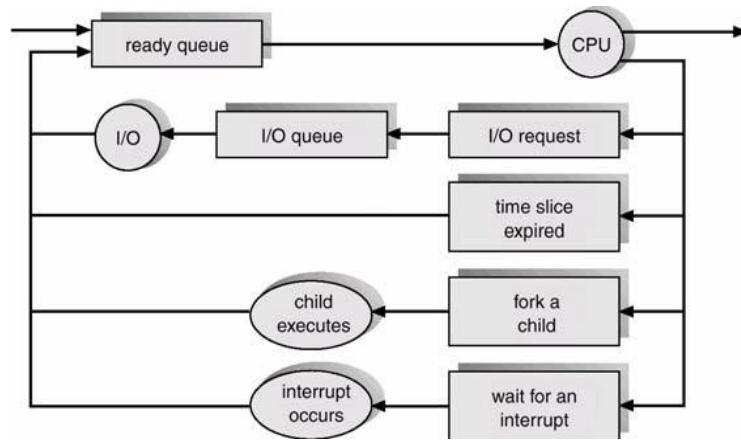
9/2/2014 (© J.P. Shen)

18-640 Lecture 3

Carnegie Mellon University 50

Process Scheduling

- **Job queue** – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- **Device queues** – set of processes waiting for an I/O device
- Process migration between the various queues



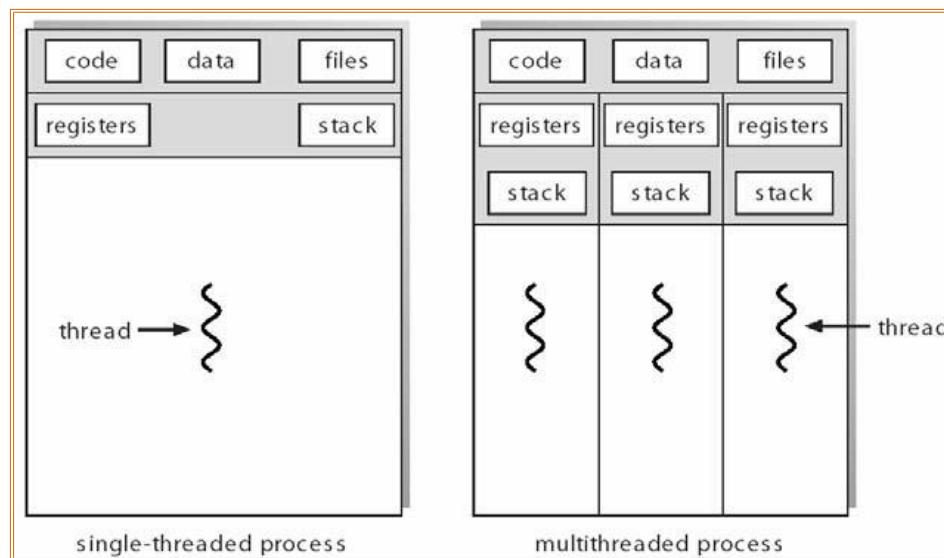
Process Creation

- Parent process creates children processes, which, in turn create other processes, forming a tree of processes
- Resource sharing
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution
 - Parent and children execute concurrently
 - Parent waits until children terminate
- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - **fork** system call creates new process
 - **exec** system call used after a **fork** to replace the process' memory space with a new program

Process Termination

- Process executes last statement and asks the operating system to decide it (**exit**)
 - Output data from child to parent (via **wait**)
 - Process' resources are de-allocated by operating system
- Parent may terminate execution of children processes (**abort**)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting
 - Some operating system do not allow child to continue if its parent terminates
 - All children terminated - *cascading termination*

Single and Multi-Threaded Processes



User Threads vs. Kernel Threads

User Threads:

- Thread management done by user-level threads library
- Three primary thread libraries:
 - POSIX Pthreads
 - Java threads
 - Win32 threads

Kernel Threads:

- Supported by the Kernel
- Examples
 - Windows XP/2000
 - Solaris
 - Linux
 - Tru64 UNIX
 - Mac OS X

Examples of Threads

- A web browser
 - One thread displays images
 - One thread retrieves data from network
- A word processor
 - One thread displays graphics
 - One thread reads keystrokes
 - One thread performs spell checking in the background
- A web server
 - One thread accepts requests
 - When a request comes in, separate thread is created to service
 - Many threads to support thousands of client requests
- RPC or RMI (Java)
 - One thread receives message
 - Message service uses another thread

Threads vs. Processes

Threads:

- A thread has no data segment or heap
- A thread cannot live on its own, it must live within a process
- There can be more than one thread in a process, the first thread calls main and has the process's stack
- Inexpensive creation
- Inexpensive context switching
- If a thread dies, its stack is reclaimed by the process

Processes:

- A process has code/data/heap and other segments
- There must be at least one thread in a process
- Threads within a process share code/data/heap, share I/O, but each has its own stack and registers
- Expensive in creation
- Expensive context switching
- If a process dies, its resources are reclaimed and all threads die

18-640 Foundations of Computer Architecture

Lecture 4: “Superscalar Implementation and Instruction Flow”

John Paul Shen

September 4, 2014

Next Time ...

- Required Reading Assignment:
 - Chapter 4 of Shen and Lipasti (SnL).
- Recommended Reference: