

18-640 Foundations of Computer Architecture

Lecture 4: “Superscalar Implementation and Instruction Flow”

John Paul Shen
September 4, 2014

- Required Reading Assignments:
- Chapters 4 and 5 of Shen and Lipasti (SnL)

- Recommended Reference:
- T. Yeh and Y. Patt, “Two-Level Adaptive Branch Prediction,” Intl. Symposium on Microarchitecture, November 1991.



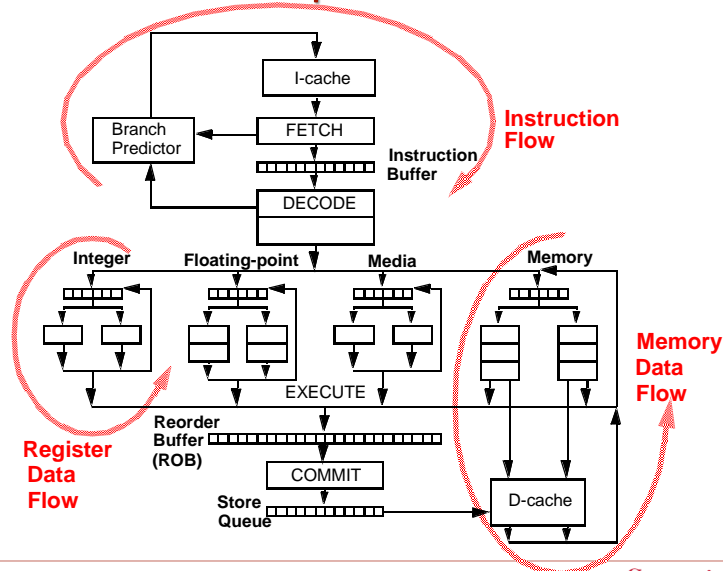
18-640 Foundations of Computer Architecture

Lecture 4: “Superscalar Implementation and Instruction Flow”

- A. Superscalar Pipeline Implementation
 - a. Instruction Fetch and Decode
 - b. Instruction Dispatch and Issue
 - c. Instruction Execute
 - d. Instruction Complete and Retire
- B. Instruction Flow Techniques
 - a. Control Flow Graph
 - b. Introduction to Branch Prediction



Three Flow Paths of Superscalar Processors

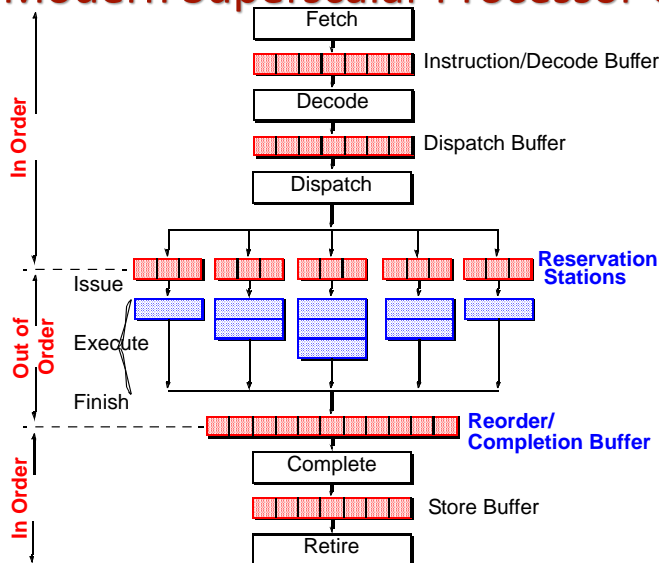


9/4/2014 (© J.P. Shen)

18-640 Lecture 4

Carnegie Mellon University 3

A Modern Superscalar Processor Organization



- Buffers provide decoupling
- In OOO designs they also facilitate (re-)ordering
- More details on specific buffers to follow

9/4/2014 (© J.P. Shen)

18-640 Lecture 4


Carnegie Mellon University 4

Superscalar Processor Implementation Issues

- **Instruction fetching Issues**
 - How do we maintain high bandwidth and accurate instruction delivery
- **Instruction decoding Issues**
- **Instruction dispatching Issues**
 - Register renaming
- **Instruction execution Issues**
 - Centralized vs distributed reservation stations
- **Instruction completion and Retiring Issues**
 - ROB, store queues, ...

A.a. Instruction Fetch and Decode

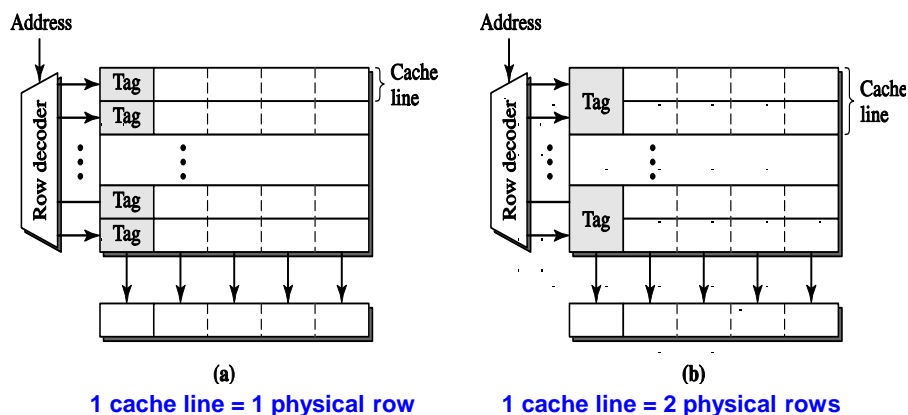
- **Goal: given a PC, fetch up to N instructions to execute**
 - Supply the pipeline with maximum number of useful instructions per cycle
 - The fetch stage **sets the maximum possible performance** (IPC_{max})
- **Impediments**
 - Instruction cache misses
 - Instruction alignment
 - Complex instruction sets
 - CISC (x86, 390, etc)
 - Branches and jumps
 - Determining the instruction address (branch direction & targets)
 - Will start in this lecture and finish in next one...



Wide Instruction Fetches

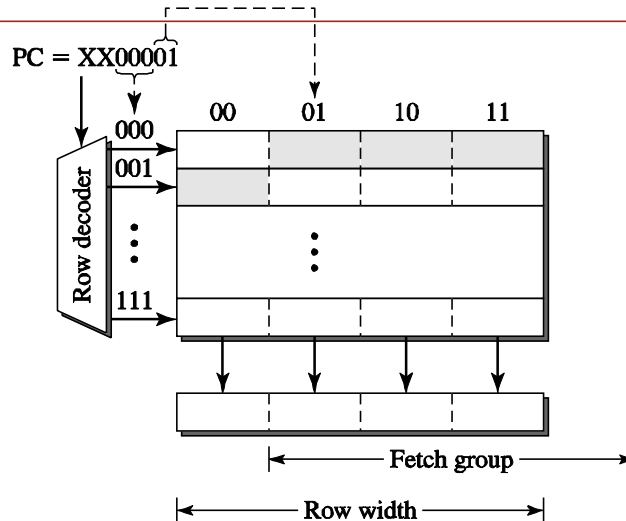
- For a N-way superscalar, need \geq N-way fetching
 - Otherwise, N-way ILP can never be achieved
 - Sometimes, wider than n-way fetch helps – why?
- Implementation: wide port to I-cache
 - Read many/all words from I-cache
 - Select those from current PC to first taken branch
- Reducing cache misses (remember?)
 - Separate I-cache
 - Larger block size, larger cache size (any problems?)
 - Higher associativity
 - 2nd-level cache, prefetching

Instruction Cache Organization



- These are logical views: In practice, tags & data may be stored separately

The Fetch Alignment Problem



9/4/2014 (© J.P. Shen)

18-640 Lecture 4

Carnegie Mellon University 9

Solving the Alignment Problem

- **Software solution:** align taken branch targets to I-cache row starts
 - Effect on code size?
 - Effect on the I-cache miss rate?
 - What happens when we go to the next chip?
- **Hardware solution**
 - Detect (mis)alignment case
 - Allow access of multiple rows (current and sequential next)
 - True multi-ported cache, over-clocked cache, multi-banked cache, ...
 - Or keep around the cache line from previous access
 - Assuming a large basic block
 - Collapse the two fetched rows into one instruction group

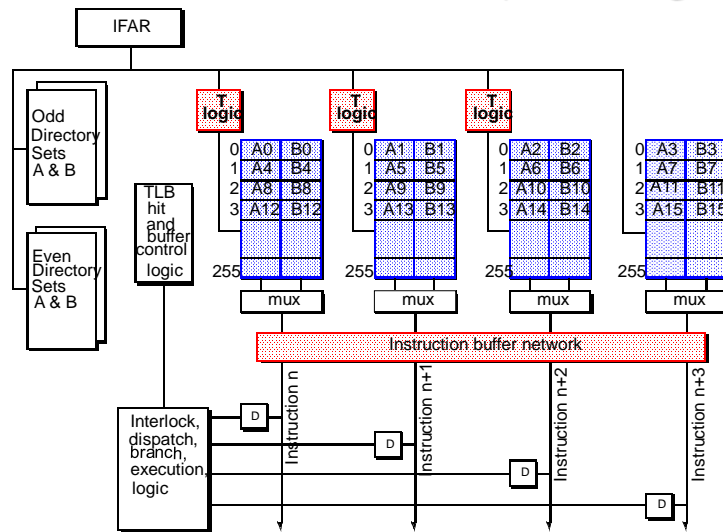
9/4/2014 (© J.P. Shen)

18-640 Lecture 4

Carnegie Mellon University 10

IBM RS/6000 I-Cache and Fetch Unit (auto alignment)

- 2-way set associative (A and B sets) I-Cache; (8) 256-instruction SRAM modules
- 16 instruction per cache line (64 bytes)



9/4/2014 (© J.P. Shen)

18-640 Lecture 4

Carnegie Mellon University 11

Cache Line Underutilization

- What if we have $< N$ instructions between two taken branches?
 - Or predicted taken branches?
- Solution: read multiple cache lines
 - Current and predicted next cache line
 - Merge instructions from two cache lines using a collapsing buffer
 - Question: how do we get the predicted next cache line address?
 - Need two predictions (PCs) per cycle
- Easier cases to handle
 - Intra-cache line forward & backward branches

9/4/2014 (© J.P. Shen)

18-640 Lecture 4

Carnegie Mellon University 12

Instruction Decoding Issues

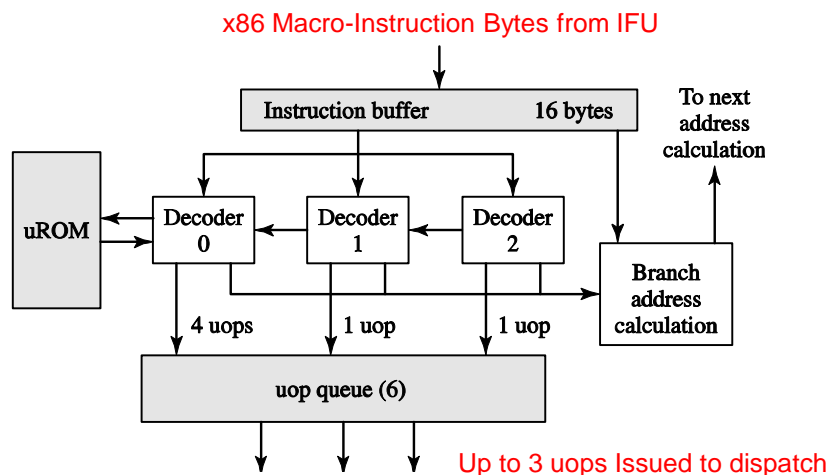
■ Primary decoding tasks

- Identify individual instructions for CISC ISAs
- Determine instruction types (especially branches)
 - Drop instructions after (predicted) taken branches
 - Potentially restart the fetch pipeline from target address
- Determine dependences between instructions

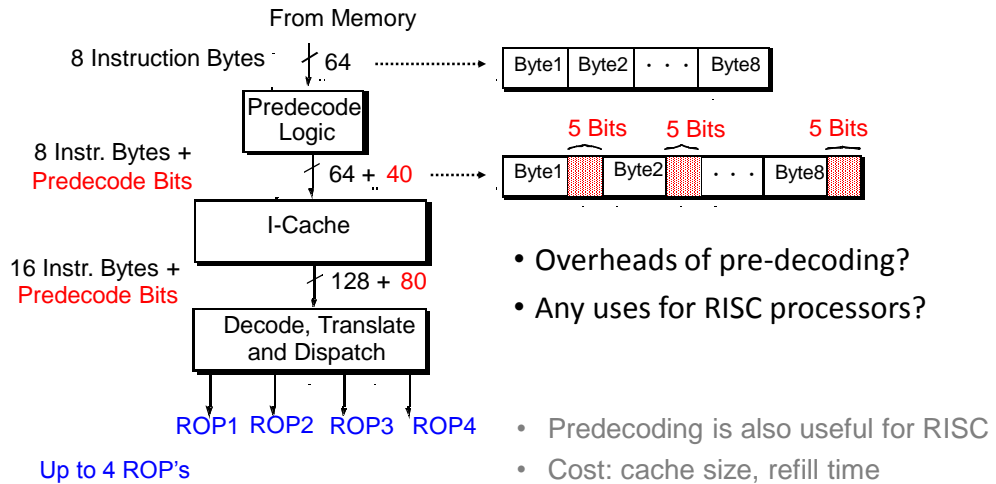
■ Two important factors

- Instruction set architecture (RISC vs. CISC)
 - Determines decoding difficulty
- Pipeline width
 - Sets the number of comparators for dependence detection

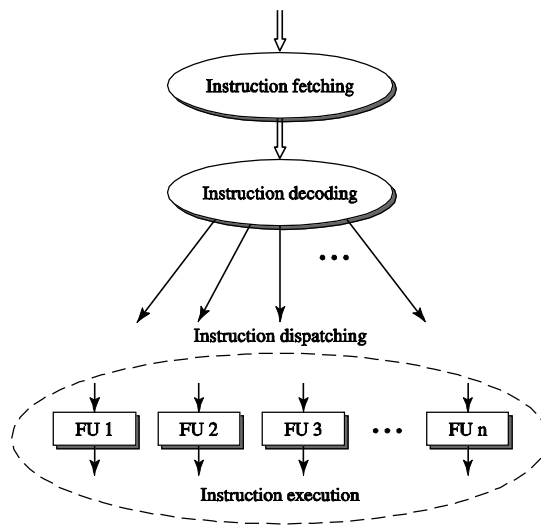
Intel Pentium Pro (P6) Fetch/Decode



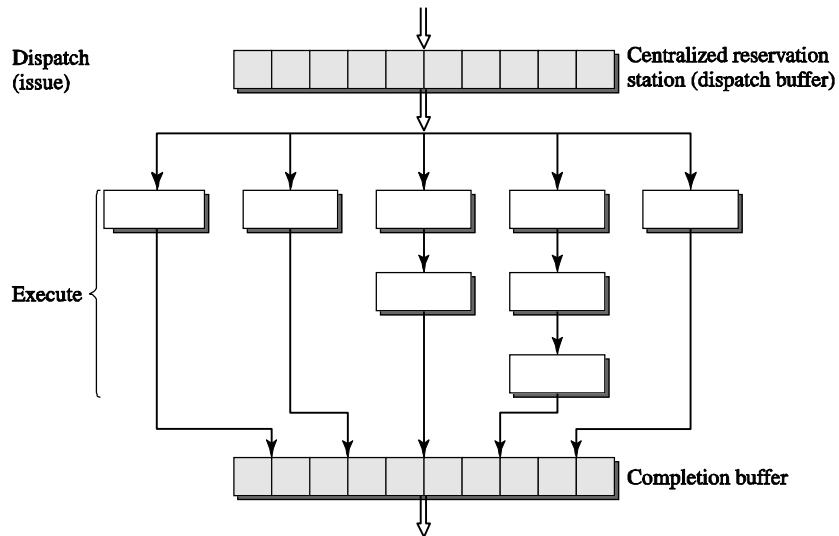
I-Cache Pre-decoding in the AMD K5



A.b. Instruction Dispatch and Issue .



Centralized Reservation Station

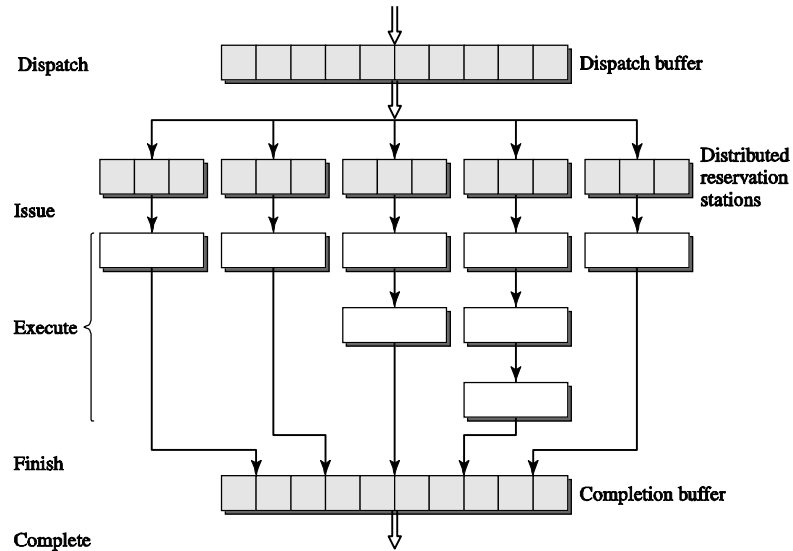


9/4/2014 (© J.P. Shen)

18-640 Lecture 4

Carnegie Mellon University 17

Distributed Reservation Stations




9/4/2014 (© J.P. Shen)

18-640 Lecture 4

Carnegie Mellon University 18

Instruction Fetch Buffer



- Smooth out the rate mismatch between fetch and execution
 - Neither the fetch bandwidth nor the execution bandwidth is consistent
- Fetch bandwidth should be higher than execution bandwidth
 - We prefer to have a stockpile of instructions in the buffer to hide cache miss latencies.
 - *This requires both raw cache bandwidth + control flow speculation*

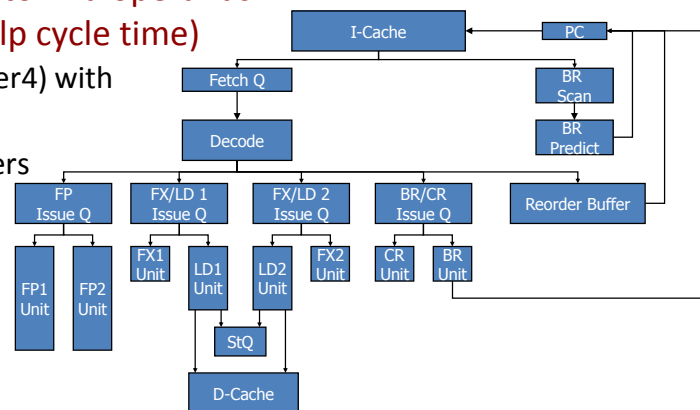
A.c. Instruction Execute

- Current trends
 - More parallelism ← forwarding/bypass very challenging
 - Deeper pipelines
 - More diversity
- Functional unit types
 - Integer
 - Floating point
 - Load/store ← most difficult to make parallel
 - Branch
 - Specialized units (media)

Forwarding/Bypass Networks

- $O(n^2)$ interconnect from/to FU inputs and outputs
- Associative tag-match to find operands
- Solutions (hurt IPC, help cycle time)

- Use RF only (IBM Power4) with no bypass network
- Decompose into clusters (Alpha 21264)




9/4/2014 (© J.P. Shen)

18-640 Lecture 4

Carnegie Mellon University 21

Specialized Functional Units

- FP multiply-accumulate
 $R = (A \times B) + C$
- Doubles FLOP/instruction
- Lose RISC instruction format symmetry:
 - 3 source operands
- Widely used




9/4/2014 (© J.P. Shen)

18-640 Lecture 4

Carnegie Mellon University 22

Specialized Functional Units in Pentium 4



- Intel Pentium 4 staggered adders – *Fireball*
- Run at 2x clock frequency
- Two 16-bit bitslices
- Dependent ops execute on half-cycle boundaries
- Full result not available until full cycle later

A.d. Instruction Complete and Retire

- **Out-of-order execution**
 - ALU instructions
 - Load/store instructions
- **In-order completion/retirement**
 - Precise exceptions
 - Memory coherence and consistency
- **Solutions**
 - Reorder buffer
 - Store buffer
 - Load queue snooping (later)

Exceptional Limitations

- **Precise exceptions: when exception occurs on instruction i**
 - Instruction i has not modified the processor state (regs, memory)
 - All older instructions have fully completed
 - No younger instruction has modified the processor state (regs, memory)
- **How do we maintain precise exception in a simple pipelined processor?**
- **What makes precise exceptions difficult on a**
 - In-order diversified pipeline?
 - Out-of-order pipeline?

Precise Exceptions and OOO Processors

- **Solution: force instructions to update processor state in-order**
 - Register and memory updates done in order
 - Usually called instruction retirement or graduation or commitment
- **Implementation: Re-Order Buffer (ROB)**
 - A FIFO for instruction tracking
 - 1 entry per instruction
 - PC, register/memory address, new value, ready, exception
- **ROB algorithm**
 - Allocate ROB entries at dispatch time in-order (FIFO)
 - Instructions update their ROB entry when they complete
 - Examine head of ROB for in-order retirement
 - If done retire, otherwise wait (this forces order)
 - If exception, flush contents of ROB, restart from instruction PC after handler

Re-Order Buffer Issues

■ Problem

- Already computed results must wait in the reorder buffer when they may be needed by other instructions which could otherwise execute.
- Data dependent instructions must wait until the result has been committed to register

■ Solution

- Forwarding from the re-order buffer
 - Allows data in ROB to be used in place of data in registers
- Forwarding implementation 1: search ROB for values when registers read
 - Only latest entry in ROB can be used
 - Many comparators, but logic is conceptually simple
- Forwarding implementation 2: use score-board to track results in ROB
 - Register scoreboard notes if latest value in ROB and # of ROB entry
 - Don't need to track FU any more; an instruction is fully identified by ROB entry

ROB Alternatives: History File

■ A FIFO operated similarly to the ROB but logs old register values

- Entry format just like ROB

■ Algorithm

- Entries allocated in-order at dispatch
- Entries updated out-of-order at completion time
 - Destination register updated immediately
 - Old value of register noted in re-order buffer
- Examine head of history file in-order
 - If no exception just de-allocate
 - If exception, reverse history file and undo all register updates before flushing

■ Advantage: no need for separate forwarding from ROB


■ Disadvantage: slower recovery from exceptions

ROB Alternatives: Future File

- Use two separate register files:
 - Architectural file \Rightarrow Represents sequential execution.
 - Future file \Rightarrow Updated immediately upon instruction execution and used as the working file.
- Algorithm
 - When instruction reaches the head of ROB, it is committed to the architectural file
 - On an exception changes are brought over from the architectural file to the future file based on which instructions are still represented in the ROB
- Advantage: no need for separate forwarding from ROB
- Disadvantage: slower recovery from exceptions

B.a. Control Flow Graph (CFG)

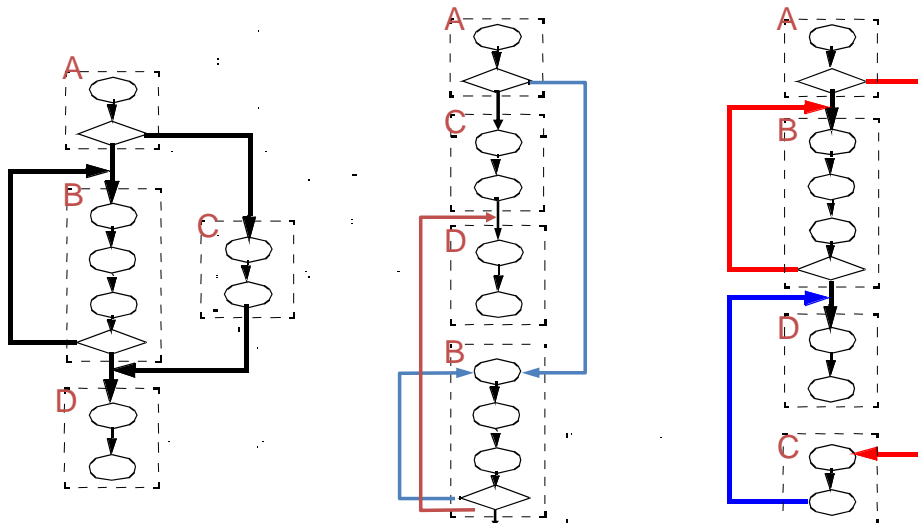
- Your program is actually a control flow graph
 - Shows possible paths of control flow through basic blocks
- Control Dependence
 - Node X is control dependent on Node Y if the computation in Y determines whether X executes



```

main:
    addi r2, r0, A
    addi r3, r0, B
    addi r4, r0, C
    addi r5, r0, N
    add r10, r0, r0
    bge r10, r5, end
loop:
    lw r20, 0(r2)
    lw r21, 0(r3)
    bge r20, r21, T1
    sw r21, 0(r4)
    b T2
T1:
    sw r20, 0(r4)
T2:
    addi r10, r10, 1
    addi r2, r2, 4
    addi r3, r3, 4
    addi r4, r4, 4
    blt r10, r5, loop
end:
  
```

Mapping CFG to Linear Instruction Sequence



9/4/2014 (© J.P. Shen)

18-640 Lecture 4

Carnegie Mellon University 31

Branch Types and Implementation

■ Types of Branches

- Conditional or Unconditional?
- Subroutine Call (aka Link), needs to save PC?
- How is the branch target computed?
 - Static Target e.g. immediate, PC-relative
 - Dynamic targets e.g. register indirect

■ Conditional Branch Architectures

- Condition Code "N-Z-C-V" e.g. PowerPC
- General Purpose Register e.g. Alpha, MIPS
- Special Purposes register e.g. Power's Loop Count

9/4/2014 (© J.P. Shen)

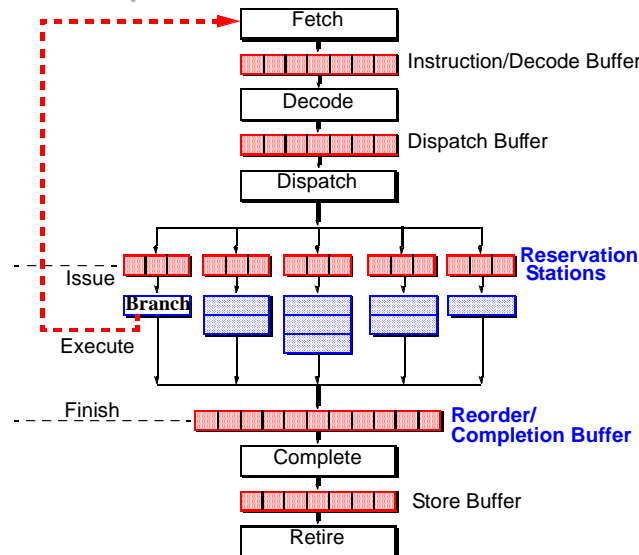
18-640 Lecture 4

Carnegie Mellon University 32

What's So Bad About Branches?

- **Robbs instruction fetch bandwidth and ILP**
 - Use up execution resources
 - Fragmentation of I-cache lines
 - Disruption of sequential control flow
 - Need to determine branch direction (conditional branches)
 - Need to determine branch target
- **Example:**
 - We have a N-way superscalar processor (N is large)
 - A branch every 5 instructions that takes 3 cycles to resolve
 - What is the effective fetch bandwidth?

Disruption of Sequential Control Flow



Riseman and Foster's Study

- 7 benchmark programs on CDC-3600
- Assume infinite machine:
 - Infinite memory and instruction stack, register file, fcn units
- *Consider only true dependency at data-flow limit*
- If bounded to single basic block, i.e. no bypassing of branches \Rightarrow maximum speedup is **1.72**
- Suppose one can bypass conditional branches and jumps (i.e. assume the actual branch path is always known such that branches do not impede instruction execution)

<i>Br. Bypassed:</i>	0	1	2	8	32	128
<i>Max Speedup:</i>	1.72	2.72	3.62	7.21	24.4	51.2

B.b. Introduction to Branch Prediction

- Why do we need branch prediction?
- What do we need to predict about branches?
- Why are branches predictable?
- What mechanisms do we need for branch prediction?

Static Branch Prediction

- **Option #1:** based on type or use of instruction
 - E.g., assume backwards branches are taken (predicting a loop)
 - Can be used as a backup even if dynamic schemes are used
- **Option #2:** compiler or profile branch prediction
 - Collect information from instrumented run(s)
 - Recompile program with branch annotations (hints) for prediction
 - See heuristics list in next slide
 - Can achieve 75% to 80% prediction accuracy
- Why would dynamic branch prediction do better?

Heuristics for Static Prediction (Ball & Larus, PPOPP1993)

Heuristic	Description
Loop Branch	If the branch target is back to the head of a loop, predict taken.
Pointer	If a branch compares a pointer with NULL, or if two pointers are compared, predict in the direction that corresponds to the pointer being not NULL, or the two pointers not being equal.
Opcode	If a branch is testing that an integer is less than zero, less than or equal to zero, or equal to a constant, predict in the direction that corresponds to the test evaluating to false.
Guard	If the operand of the branch instruction is a register that gets used before being redefined in the successor block, predict that the branch goes to the successor block.
Loop Exit	If a branch occurs inside a loop, and neither of the targets is the loop head, then predict that the branch does not go to the successor that is the loop exit.
Loop Header	Predict that the successor block of a branch that is a loop header or a loop preheader is taken.
Call	If a successor block contains a subroutine call, predict that the branch goes to that successor block.
Store	If a successor block contains a store instruction, predict that the branch does not go to that successor block.
Return	If a successor block contains a return from subroutine instruction, predict that the branch does not go to that successor block.

Dynamic Branch Prediction Tasks

• Target Address Generation

- Access register
 - PC, GP register, Link register
- Perform calculation
 - +/- offset, auto incrementing/decrementing

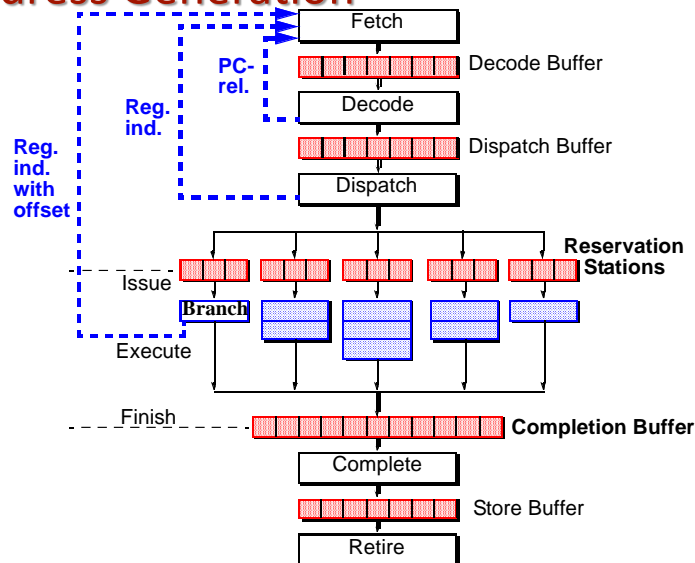
⇒ Target Speculation

• Condition Resolution

- Access register
 - Condition code register, data register, count register
- Perform calculation
 - Comparison of data register(s)

⇒ Condition Speculation

Target Address Generation



Determining Branch Target

Problem: Cannot fetch subsequent instructions until branch target is determined

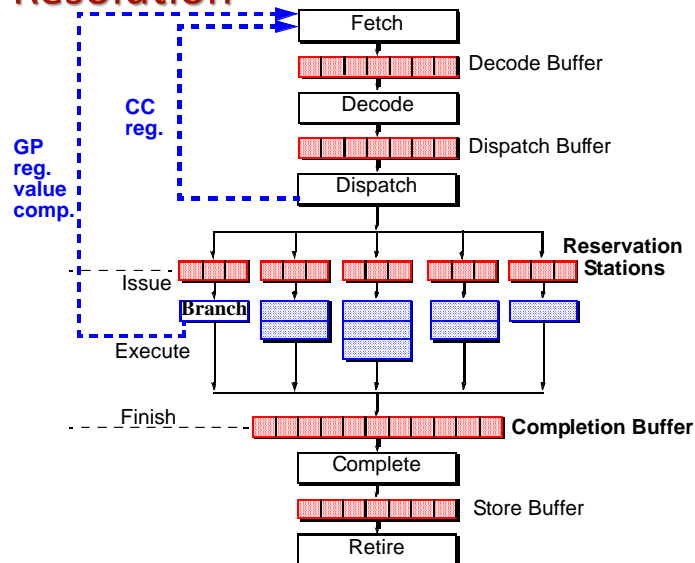
■ Minimize delay

- Generate branch target early in the pipeline

■ Make use of delay

- Bias for not taken
- Predict branch targets
- For both PC-relative vs register Indirect targets

Condition Resolution



Determining Branch Direction

Problem: Cannot fetch subsequent instructions until branch direction is determined

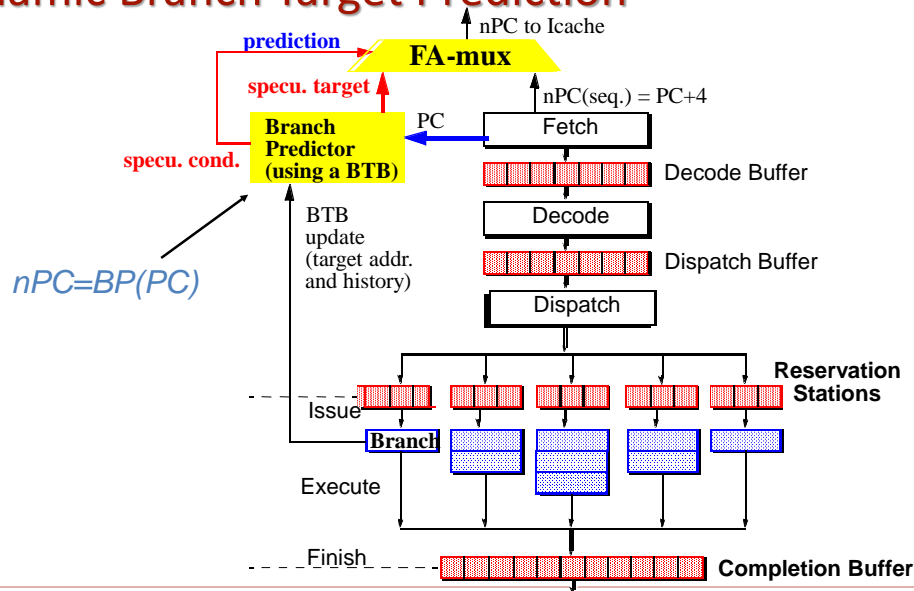
■ Minimize penalty

- Move the instruction that computes the branch condition away from branch (ISA & compiler)
 - 3 branch components can be separated
 - Specify end of BB, specify condition, specify target


■ Make use of penalty

- Bias for not-taken
- Fill delay slots with useful/safe instructions (ISA & compiler)
- Follow both paths of execution (hardware)
- *Predict branch direction (hardware)*

Dynamic Branch Target Prediction



Target Prediction: Branch Target Buffer (BTB)



- A small “cache-like” memory in the instruction fetch stage
- Remembers previously executed branches, their addresses (PC), information to aid target prediction, and most recent target addresses
- I-fetch stage compares current PC against those in BTB to “guess” nPC
 - If matched then prediction is made else $nPC = PC + 4$
 - If predict taken then $nPC = \text{target address in BTB}$ else $nPC = PC + 4$
- When branch is actually resolved, BTB is updated

More on BTB (aka BTAC)

- Typically a large associative structure
 - Pentium3: 512 entries, 4-way; Opteron: 2K entries, 4-way
- Entry format
 - Valid bit, address tag (PC), target address, fall-through BB address (length of BB), branch type info, branch direction prediction
- BTB provides both target and direction prediction
- Multi-cycle BTB access?
 - The case in many modern processors (2 cycle BTB)
 - Start BTB access along with I-cache in cycle 0
 - In cycle 1, fetch from BTB+N (predict not-taken)
 - In cycle 2, use BTB output to verify
 - 1 cycle fetch bubble if branch was taken

Branch Condition Prediction

- **Biased For Not Taken**
 - Does not affect the instruction set architecture
 - Not effective in loops
- **Software Prediction**
 - Encode an extra bit in the branch instruction
 - Predict not taken: set bit to 0
 - Predict taken: set bit to 1
 - Bit set by compiler or user; can use profiling
 - Static prediction, same behavior every time
- **Prediction Based on Branch Offsets**
 - Positive offset: predict not taken
 - Negative offset: predict taken
- **Prediction Based on History**

9/4/2014 (© J.P. Shen)

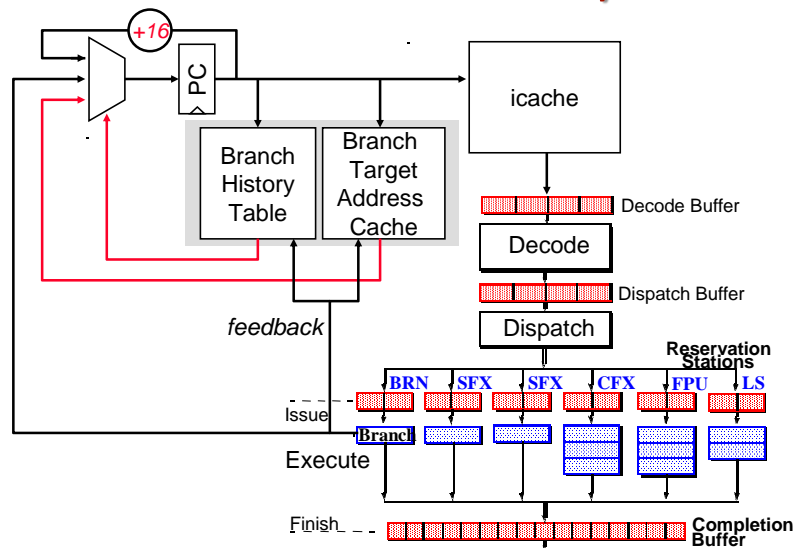
18-640 Lecture 4

Carnegie Mellon University 47

Dynamic Branch Prediction Based on History

- Use HW tables to track history of direction/targets
 - $\text{nextPC} = \text{function}(\text{PC}, \text{history})$
- Need to verify prediction
 - Branch still gets to execute

More Next Time....



9/4/2014 (© J.P. Shen)

18-640 Lecture 4

Carnegie Mellon University 48