

# 18-640 Foundations of Computer Architecture

## Lecture 11: “Dynamic Code Translation & Compilation”

Mauricio Breternitz, Ph.D., AMD Fellow  
October 2, 2014

➤ Recommended Reading Assignment:

- “Dynamic Binary Modification: Tools, Techniques, and Applications”  
by Kim Hazelwood, in *Synthesis Lectures on Computer Architecture*,  
Morgan & Claypool Publishers, 2011.



# 18-640 Foundations of Computer Architecture

## Lecture 11: “Dynamic Code Translation & Compilation”

- A. Dynamic Binary Translation
- B. Transmeta Example
- C. Code Morphing System
- D. Nvidia Denver Example



## A. DBT: Dynamic Binary Translation

- Up to now, CPU hardware design
  - Pipelining, superscalar, speculative execution, ...
  - ISA to interface hardware features (VLIW)
- For VLIW, the compiler's role has been static
  - Hence, it's effectiveness limited
- DBT provide the compiler with a dynamic role

## What is Binary Translation?

- Translating programs in one binary format to another
  - Different ISAs
    - E.g. PowerPC => x86 : to port programs across platform
  - Same ISAs
    - E.g. x86 => x86 : code optimization or feature instrumentation
  - Intermediate Representation
    - E.g. Java bytecode => x86 : to avoid interpretation
- When
  - Static : translation before running programs
  - Dynamic : translation while running programs

## Why is DBT a Good Idea?

- Feature support without hardware/source-code modification
  - E.g. Binary compatibility, optimization
- Vs. static binary translation
  - Access to complete program
    - Programs are fully linked
  - Access to program state
    - Include dynamic values
  - Can handle self-modifying code
  - Can **adapt to changes** in program behavior
  - No need to re-link
    - Translate instruction and jump to it

## Dynamic Compilation Challenges

- Managing compilation costs
  - Trade-off compilation speed for generated code quality
  - Dynamic and static memory footprint important in certain systems
- Binary sources have less information than original source
  - Decompilation of CFG/DFG from binary
  - Maintain exception behavior of original code


## Dynamic Compilation Challenges (cont.)

- Integrating many components together
  - Compilers
  - Profilers
- On-stack replacement
  - Application
    - Dynamic optimization of procedures on the call stack
    - Mid-iteration optimization of long-running loops
  - Transition from old code to new compiled code
    - Stack frames of recompiled methods must match original if they are currently on the runtime stack
    - Optimized code may use stack in strange ways
    - Must be able to perform from inlined context

## Managing Compilation Costs

- Mixed-mode / lazy compilation
  - Apply expensive optimizations to frequently executed code
  - Possible combinations
    - Combine interpreter / optimizing compiler
    - Combine fast & optimizing compiler
- Use less aggressive compiler optimizations
  - Avoid expensive/iterative global optimizations
    - Local optimizations have larger payoff
  - Use faster register allocation schemes
    - Linear scan vs. graph coloring
- Implement efficient compiler
  - Apply optimizations concurrently
  - Pipeline dataflow information between stages



## Dynamic Binary Translation Overview



- Dynamic program modifier
  - Start with interpretation or intercept program execution
  - Observe a sequence of instructions
  - Produce new code and save in “code cache”
    - Change the jump/branch target address properly
  - Manipulate or add instructions as needed



## DBT Configurations (1)

- Cross platform
  - E.g. Crusoe (Transmeta)
- Same platform
  - E.g. Dynamo (HP), PIN (x86), Dynamo-Rio(x86)

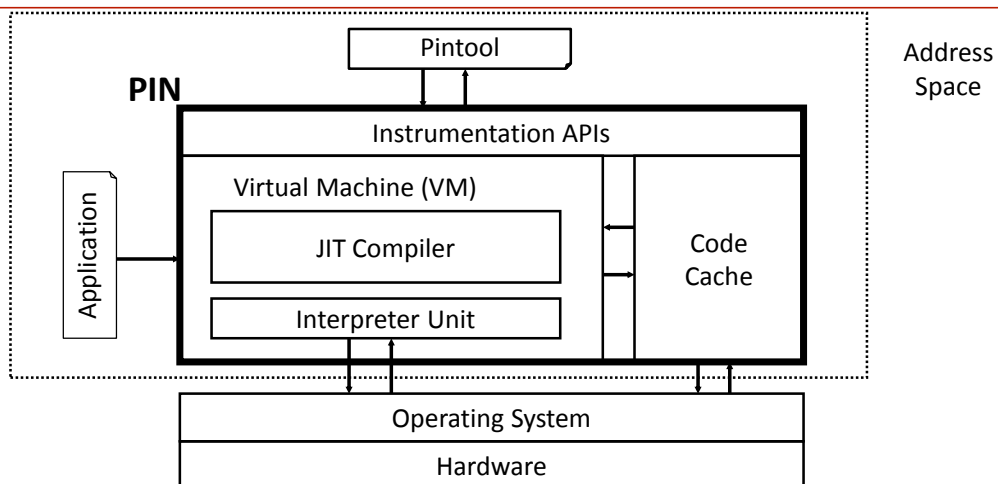


## DBT Configurations (2)

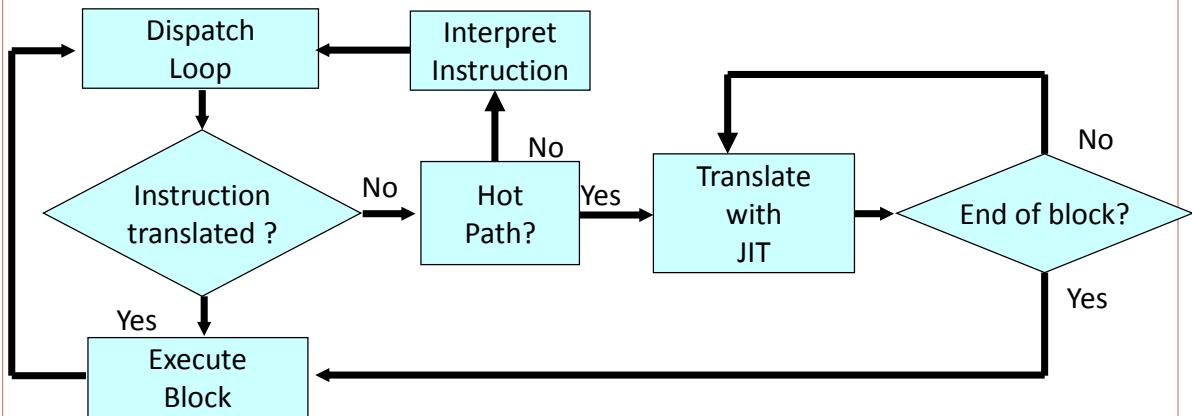
- Virtual Machine
  - E.g. ESX server (vmWare)
- JIT compilation
  - E.g. JVM (Sun), C#(MS)



## Example DBT : PIN for x86



## Translation Loop in VM



\* Some DBTs have no interpretation mode

10/2/2014 (© J.P. Shen)

18-640 Lecture 11

Carnegie Mellon University 13

## Code Cache

- Software cache in virtual address space
  - Keep translated code in memory for later reuse
  - Essential to leverage the high cost of translation and optimization
  - Trade-off: cost of memory vs. higher reuse
- Allocation policy
  - Remember every translated blocks: doubling code size at least
  - Pick up hot path: temporal locality
- Granularity
  - Basic block (easy but frequent switches to interpretation)
  - Trace (better amortization of switch overhead but need path profiling)

10/2/2014 (© J.P. Shen)

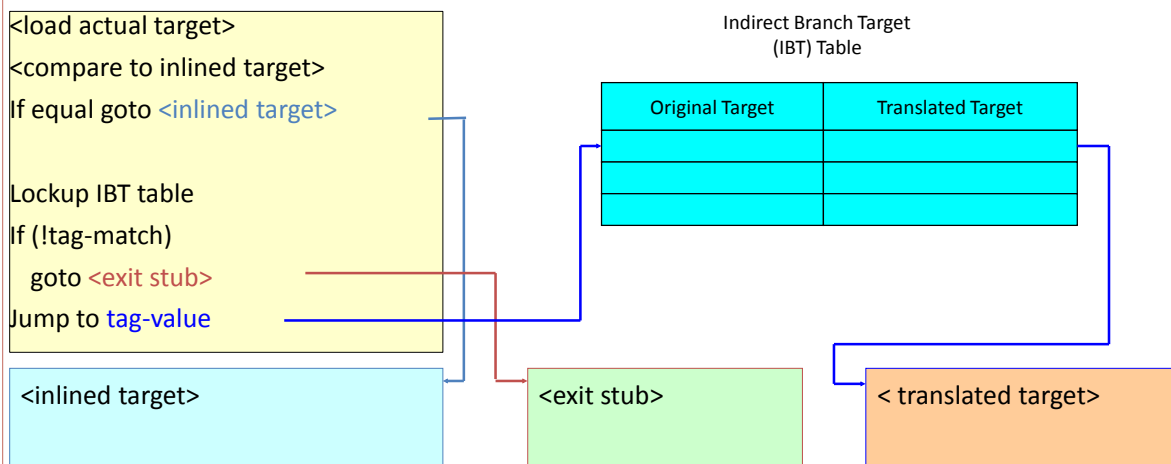
18-640 Lecture 11

Carnegie Mellon University 14

## Linking Translated Basic Blocks (1)

- Context-switch overhead between interpretation and translation
  - Additional instructions for context switch
  - Register Spilling
  - Limited DBT code optimization boundary
- Link translated code
  - Direct branches are easy (branch to proper code cache offset)
  - Indirect branch target (IBT) table for indirect branch
    - (key : value) = (original target, translated target)
- Inline preferred target
  - Inline the block of the preferred target
  - Add check code

## Linking Translated Basic Blocks (2)





## Exceptions

- Asynchronous exceptions (interrupts)
  - Can be delayed, easy
  - Wait until the current translated code finishes
  - Translate exception handler
  - Invoke translated exception handler
- Synchronous exception
  - During interpretation, no problem
  - During executing translated code, either revert the instruction execution and interpret
    - E.g. checkpoint support in Crusoe's VLIW hardware
  - Or make sure to stop at exact point
    - E.g. NullPointerException in Java
  - Invoke translated exception handler

## Self-Modifying Code

- Solution 1 : stick to interpretation
  - Modifying code located in heap
  - Always interpret when jumping to heap address
- Solution 2 : invalidate translated code
  - When jumping to heap address, translate the code and cache it
  - Write-protect the pages with modifying code
  - Execute translated code
  - If the code changes later, page-fault exception is triggered
  - Invalidate corresponding translated code of the page in code cache
  - Next time, the code is re-translated

## Multicore – Access to shared data

- Speculative execution assumes data unchanged
- Access by another core invalidates this assumption
- Solution:
  - “Victim” cache stores speculative state
  - Access to shared data aborts “commit” of translation

## Example Use for DBT (Dynamo)

- Dynamic optimization by software
  - Dynamic : leverage runtime information (compared with compiler)
  - Software : flexible, sophisticated (compared with out-of-order execution)
    - But, with time constraints
- Runtime information
  - Trace, DLL, function call counter, dynamic values
- Optimizations
  - Superblock scheduling: allows additional chance for classic optimization
    - ILP scheduling, copy propagation, loop unrolling
  - Inlining : balance code size and function call overhead
  - Fast shared library invocation : remove lookup overhead
  - Register allocation: reduce register spill

## Other DBT Uses

- Profiling
- Architectural studies
- Security/debugging tools
- ...
- Most of this uses require instrumentation API
  - Define uses through the API by defining which instructions are examine and what code is inserted

## DBT Used for Binary Instrumentation

- MIPS pixie/pixstats ~1987
  - Statically modified binaries to count instructions
- Sun Spix/Spixstats/Span/Shadow/Shade ~1988-95
  - Success led to more functionality
  - Ability to patch in analysis routines while running
  - Shade ~1990 enhanced dynamic techniques
    - able to run MIPS and x86 code on SPARC, about 1/3 native speed
    - Amazed at the performance of non-native code
- Lessons –Critical realizations in 1995
  - Co-design the underlying hardware to reduce inefficiencies
  - Could move from less than native to net performance gain
  - Meant Hybrid processors could compete with mainstream

## B. Transmeta – A Real-Life DBT system

- Transmeta
- Hardware Support for Binary Translation
- DBT Optimizations

## Binary Translation Uses

- Pentium Pro and its successors
  - Translates x86 binary code into internal UOPS via HW
- Intel IA32EL
  - Translates user level x86 programs on Itanium via SW
  - Performance about 60% of native machine
- JAVA JITs / Microsoft MSIL / VMware
- Transitive Technologies
  - Rosetta runs Apple PowerPC user programs on x86
- Transmeta Crusoe and Efficeon mobile processors
  - Translates x86 binary code into internal UOPS via sw
  - Transparent full system level translation as real products
- Performance comparable to Intel's mobile CPUs

## Binary Translation History

- Good work worth mentioning:

- |                      |                          |
|----------------------|--------------------------|
| • XDOS               | Early static DOS to Unix |
| • DEC FX!32          | x86 user apps to Alpha   |
| • IBM DAISY, BOA     | PowerPC on VLIW          |
| • Elbrus nArch, E2K  | SPARC, x86 to VLIW       |
| • Sun SoftWindows    | x86 to SPARC/Solaris     |
| • Intel PINx86       | binary instrumentation   |
| • Intel StarDBTIntel | research translator      |

and many others ...

- see Virtual Machines book by Jim Smith & Ravi Nair

## Transmeta “Hybrid Processor”

- A Hybrid processor –in Transmeta’s proposal

Specifically means a microprocessor implemented as a  
hardware/software co-design using software binary translation and  
optimization with special purpose hardware

so that

it can function as a fully compatible microprocessor with performance  
comparable to pure hardware implementations.

## Transmeta Crusoe Characteristics

- Transmeta's Crusoe microprocessor is a full, system level implementation of the x86 architecture, comprising a native VLIW microprocessor with a software layer, the Code Morphing Software (CMS)
  - 128bit VLIW(very long instruction word) engine
    - The Crusoe processor features a 128-bit wide VLIW (Very Long Instruction Word) engine that can issue up to 4 instructions per clock cycle.
  - Code Morphing Software
    - Code Morphing Software (CMS) layer provides the Transmeta Crusoe processor with x86 compatibility
  - Integrated Architecture
    - To ease system design and enhance performance, Transmeta has integrated Northbridge functionality — SDR and DDR SDRAM memory controllers, a 32-bit, 33MHz PCI bus controller and a Serial ROM interface controller — directly into the Crusoe processor die.
  - LongRun Technology
    - Transmeta LongRun technology allows the Transmeta Crusoe processor to conserve power by dynamically adjusting its voltage and clock frequency.

## Transmeta

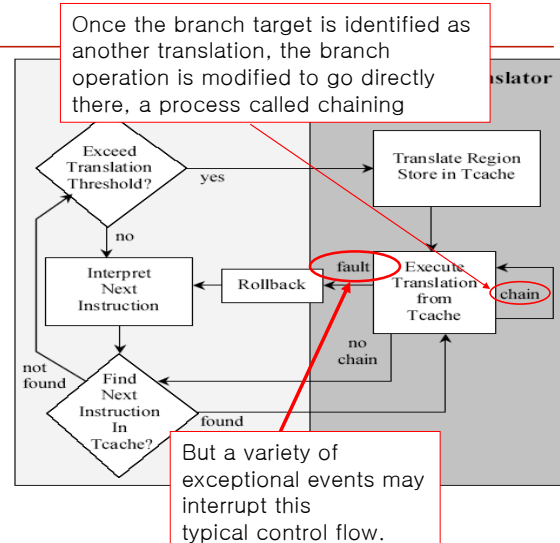
- Transmeta was founded in 1995, with the idea to co-design hardware and software together in a machine whose goal was provide an efficient base for binary translation of x86 programs
- Spent \$600M in R&D over 12 years to learn the tricks
- CMS was invisible to users, no compatibility bugs
- Five generations of processors designed
  - First two never shipped, insufficient performance
  - Crusoe product—Pentium class for mobile (250nm-180nm)
  - Efficeon product—Pentium-M class for mobile (130nm-90nm)
  - Next generation design that never shipped
- Lesson: Worked well and was on a rapid learning curve
- IPC improving at roughly 2x / generation

## C. Transmeta CMS – “Code Morphing System”

- To produce high performance while remaining perfectly faithful to an existing architecture, translator must optimize aggressively:
  - Speculation: Translator makes aggressive assumptions about code to achieve higher performance
  - Recovery: Check assumptions and rollback to commit points if they prove to be false, for precise interpretation
  - Adaptive retranslation: If recovery is required too often, retranslate with less aggressive assumptions

## Overview of CMS Control Flow

- Typical CMS control flow
  - CMS is structured like many dynamic translation systems.
  - Initially, an interpreter decodes and executes x86 instructions sequentially.
  - When the number of executions of a section of x86 code reaches a certain threshold, its address is passed to the translator.
  - The translator selects a region and stores the translation with various related information in the translation cache.
  - From then on, until an event invalidates the translation cache entry, CMS executes the translation when the x86 flow of control reaches the translated code region.



## Hardware Support for Recovery

- *Shadow registers*:
  - *Working* and *shadow* copies of x86 registers
- *Commit atom*: Copies working registers to shadow registers, commits memory writes
- *Rollback atom*: Copies shadow registers to working registers, discards memory writes

## Example: Precise Exceptions

- **Problem**: CMS rearranges operations in a translation, but x86 has precise exception semantics
- **Speculation**: CMS translations scheduled assuming no exceptions will occur
- **Recovery**: If an exception occurs, rollback to preceding consistent commit point, and interpret sequentially
- **Adaptive retranslation**: An instruction causing exceptions too often is isolated, and the rest of the original translated code is retranslated so it won't need rollback




## Memory-Mapped I/O Problem

- **Problem:** Arbitrary x86 memory operations may be memory-mapped I/O, which must not be reordered and may not be rolled back
- **Speculation:** CMS translates memory operations as *normal*, reorders them at will (unless marked by interpreter as *abnormal*)
- **Recovery:** A normal reference to abnormal memory causes a hardware fault, rollback, and appropriate interpretation
- **Adaptive retranslation:** A memory op that is abnormal too often is retranslated with commits before and after it, and with other constraints necessary for memory-mapped I/O

## Memory-Mapped I/O

Figure 2: Degradation Caused by Suppressing Memory Reordering



## Data Speculation (Aliases)


- **Problem:** To achieve effect of out-of-order processors on memory operations, CMS needs to reorder memory operations
- **Speculation:** CMS assumes memory operations don't alias unless it can prove otherwise, reorders accordingly
- **Recovery:** Memops scheduled earlier set an *alias register*; later memops that might alias check the register and trap on overlap
- **Adaptive retranslation:** Translation that takes alias faults too often is translated without reordering

## Data Speculation

Simulations obtained  
without using alias HW

No adaptive retranslation  
cuts PCmark2002 3d vector  
performance by a factor of  
47.5 on hardware

Figure 3: Degradation Caused By No Alias Hardware



## SMC – Self-Modifying Code

- **Original problem:** If the x86 code modifies itself, the CMS translations must be invalidated or otherwise adapt.
- **Speculation:** Normal translations assume no SMC
- **Simple recovery:** Write-protect x86 code pages, find and invalidate corresponding translations if a fault occurs
- **Secondary problems:**
  - –Inefficient for self-modifying code: granularity too large
  - –Can't distinguish data in same page as code
- **Costs incurred by CMS:**
  - –Handling fault, invalidating translations, special processing
  - –Generating new translations for new code

## SMC- Fine-Grain Protection

- **First refinement:** Hardware support for sub-page granularity  
Only needed for a few pages at a time, allowing tiny hardware cache
- Without fine-grain protection, the worst cases:

	Faults	Slowdown
Win95 Boot	52.8x	2.2x
Win98 Boot	59.4x	3.8x
MultimediaMark	46.8x	1.6x
WinStone Corel	54.2x	2.1x
Quake Demo2	7.7x	1.02x

## Transmeta – 10 Hardware Support Features

1. Software controlled state. Commit/Rollback/Abort
2. More registers than architected state
3. Alias detection under software control
4. Self modifying code detection with fine grain support
5. Auto-typing of pure memory vs I/O
6. Fast traps supported by underlying runtime system
7. Instruction primitives for fast interpretation
8. Private memory. ~5% of DRAM for translated code
9. Private non-volatile storage (FLASH ROM) with every chip
10. Competitive CPU hardware, ie uOp ISA, IPC and clock rate

## Speculative Execution and x86 State Control

- Hardware for software controlled Commit/Rollback enables wide optimization regions

All registered state has two copies

- Working state
- Committed State

Registered State includes:


- Integer and Floating Point register files
- Condition code flags
- All other state registers

Support for speculative stores between commit points in L1 data cache

- New cache state
- 1 new tag bit “speculative”
- Flash (1-cycle) flush (on rollback) or convert to dirty (commit)

- New UOP level Instructions

- Commit: Copies working registers to committed state (1-cycle)
- Rollback: Copies committed registers to working registers (1-cycle)
- Abort: To runtime trap handler, which can do fix up, rollback, and branch



## Code Morphing Software

- 4 “Gear” System
  - Start with interpreter
  - Increasing levels of optimization and speculation
  - Utilize HW assistance
    - Speculation
      - Shadow Registers
      - Speculative data in cache

### First Gear



**1<sup>st</sup> Gear**


Executes 1 instruction at a time

- Profiles code at runtime
- Gathers data for flow analysis
- Gathers branch frequencies and directions
- Detects load/store typing (IO vs memory)

Filters out infrequently executed code

**No startup cost  
Lowest speed**

## Second Gear




Uses profile data to create initial translations after code reaches 1<sup>st</sup> threshold.

- Translates a "Region" of up to 100 x86 instructions.
- Adds flow graph "Shape" information
- Light Optimization
- "Greedy" scheduling

Low translation overhead  
Fast execution

## Third Gear




Further optimizes the 2<sup>nd</sup> gear regions

- Common sub-expression elimination
- Memory re-ordering
- Significant code optimization
- Critical path scheduling

Medium translation overhead  
Faster execution

## Fourth Gear



Most advanced optimizations for "hottest" code regions.

- Splices together multiple regions
- Optimizes across region boundaries
- Used advanced behavioral data
- Critical path scheduling

Highest translation overhead  
Fastest execution

## CMS Optimizations

- Aggressive scheduling of instruction level parallelism
- Out-of-Order instruction execution obtained on In-Order hardware
- Critical path height reduction
- Common sub-expression elimination
- Uses "Address Alias Checking Hardware"
  - Re-ordering of loads and stores even with potential aliases
  - Elimination of loads and stores even with potential aliases
- Software register renaming
- Fusing operations
- Dead code elimination
- Removal of conditional branches
- Adaptive re-translation during program execution
- Loop unrolling and optimization
  - Remove Exit Branches      Loop invariant code motion
  - Code motion across back-edge      Strength reduction

HOT CODES  
GET  
MOST  
OPTIMIZATION

## Translation Example

### translation example

#### → X86 instruction

A. `addl %eax, (%esp)` // load data from stack, add to %eax  
 B. `addl %ebx, (%esp)` // ditto, for %ebx  
 C. `movl %esi, (%ebp)` // load %esi from memory  
 D. `subl %ecx, 5` // subtract 5 from %ecx register

#### → In a first pass, the front end of the translation – simple translation


```
ld %r30, [%esp] // load from stack, into temporary
add.c %eax, %eax, %r30 // add to %eax, set condition codes.
ld %r31, [%esp]
add.c %ebx, %ebx, %r31
ld %esi, [%ebp]
sub.c %ecx, %ecx, 5
```

#### → In a second pass, the optimizer. applying well known compiler optimization skill such as common subexpression elimination, loop invariant removal or dead code elimination.

```
ld %r30, [%esp] // load from stack only once
add %eax, %eax, %r30
add %ebx, %ebx, %r30 // reuse data loaded earlier
ld %esi, [%ebp]
sub.c %ecx, %ecx, 5 // only this last condition code needed
```

#### → In a final pass, the scheduler. reordering atoms into molecules.

```
1. ld %r30, [%esp]; sub.c %ecx, %ecx, 5
2. ld %esi, [%ebp]; add %eax, %eax, %r30; add %ebx, %ebx, %r30
```



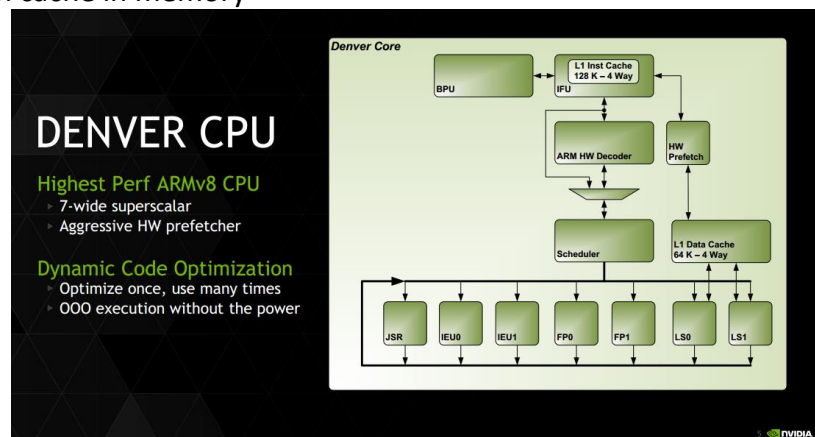


## Code Morphing Software

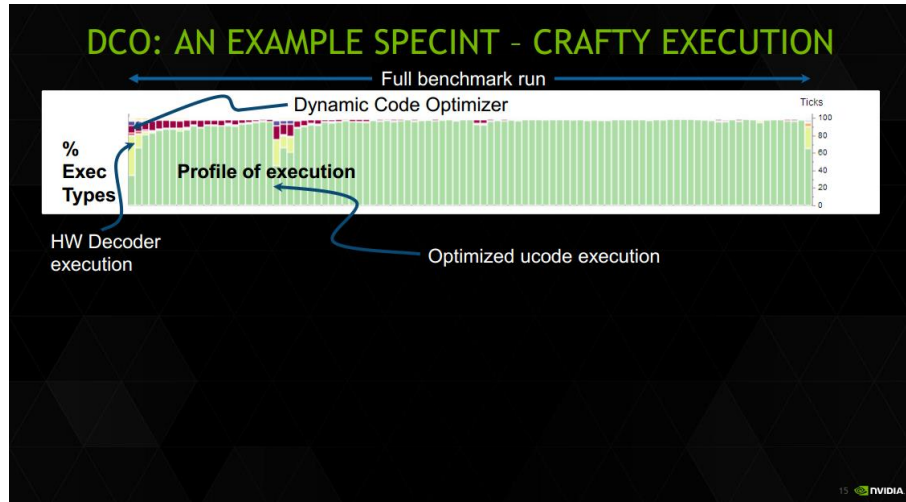
- Multicores
- Omission from Initial Transmeta proposal: multicores
  - Coherence
    - Rollback on “speculative” data hit
    - Research at Intel

## D. NVIDIA “Denver” Architecture

- 128 Mb translation cache in memory
- VLIW engine



## NVIDIA Denver



10/2/2014 (© J.P. Shen)

18-640 Lecture 11

Carnegie Mellon University 51

## Binary Translation Advantages

- Innovation
  - To allow processor innovation not tied to particular instruction sets
  - Using BT to provide backwards compatibility
- Performance
  - To enable new means to improve processor performance
  - Using BT to provide backwards compatibility
- Power
  - To enable simpler and lower power processors
  - Using BT to provide backwards compatibility

10/2/2014 (© J.P. Shen)

18-640 Lecture 11

Carnegie Mellon University 52

## Conclusion

- Binary Translation based processors can work well, no longer a theory.
- Special purpose hardware support is needed, co-designed with software.
- Software translation can be done poorly or well, special care is needed to keep translation overhead low.
- Many opportunities for clever hardware/software co-design tradeoffs
- This is a technological approach still in its infancy

## References

- Wang, Cheng, et al. "Stardbt: An efficient multi-platform dynamic binary translation system." *Advances in Computer Systems Architecture*. Springer Berlin Heidelberg, 2007. 4-15.
- Chernoff, Anton, et al. "FX! 32: A profile-directed binary translator." *IEEE Micro* 18.2 (1998): 56-64.
- Cmelik, Bob, and David Keppel. *Shade: A fast instruction-set simulator for execution profiling*. Springer US, 1995.
- Dehnert, James C., et al. "The Transmeta Code Morphing™ Software: using speculation, recovery, and adaptive retranslation to address real-life challenges." *Proceedings of the international symposium on Code generation and optimization: feedback-directed and runtime optimization*. IEEE Computer Society, 2003.
- Sites, Richard L., et al. "Binary translation." *Communications of the ACM* 36.2 (1993): 69-81.
- Ebcioglu, Kemal, et al. "Dynamic binary translation and optimization." *Computers, IEEE Transactions on* 50.6 (2001): 529-548.