

18-640 Foundations of Computer Architecture

Lecture 14: “Multicore Cache Coherence”

John Paul Shen
October 23, 2014

➤ Required Reading Assignment:

- “Parallel Computer Organization and Design,” by Michel Dubois, Murali Annavaram, Per Stenstrom, Chapters 5 and 7, 2012.



Carnegie Mellon University ¹

10/23/2014 (© J.P. Shen)

18-640 Lecture 14

18-640 Foundations of Computer Architecture

Lecture 14: “Multicore Cache Coherence”

- A. Cache Coherence Problem
- B. Cache Coherence Protocols
 - Write Update
 - Write Invalidate
- C. Bus-Based Snoopy Protocols
 - VI & MI Protocols
 - MSI, MESI, MOESI Protocols
- D. Directory-Based Protocols



Carnegie Mellon University ²


10/23/2014 (© J.P. Shen)

18-640 Lecture 14

Shared Memory Multiprocessors


- All processor cores have access to unified physical memory
 - They can (implicitly) communicate using loads and stores
- Advantages
 - Looks like a better multithreaded processor (multitasking)
 - Requires evolutionary changes to the OS
 - Threads within an app communicate implicitly without using OS
 - Simpler to code for and lower overhead
 - App development: first focus on correctness, then on performance
- Disadvantages
 - Implicit communication is hard to optimize
 - Synchronization can get tricky
 - Higher hardware complexity for cache management

Shared Memory Multiprocessors




- Caches are (equally) helpful with multiprocessors
 - Reduce access latency, reduce bandwidth requirements
 - For both private and shared data across processors
- But caches introduce the problems of coherence & consistency

Cache Coherence Problem: Example



- Processors may see different values for *u* after event 3
- With write back caches, value written back to memory depends on which cache flushes or writes back value when doing write back
 - Threads or processes accessing main memory may see very stale value
- Unacceptable for programming, and it's frequent!

Hardware Cache Coherence Using Snooping



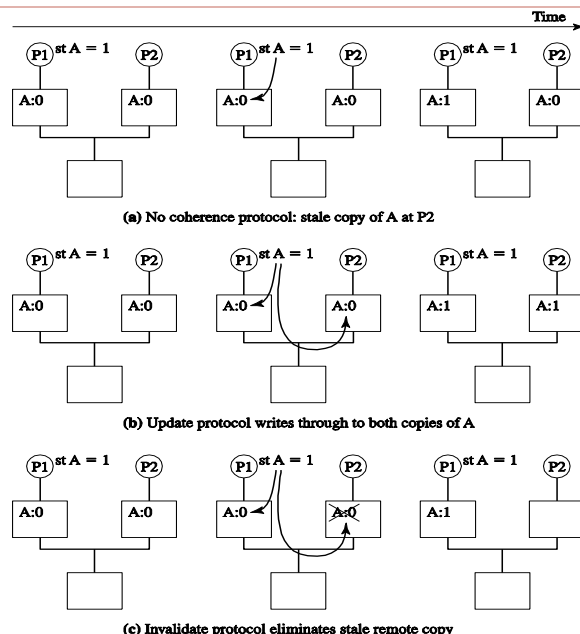
- Hardware guarantees that loads from all cores will return the value of the latest store
- Coherence mechanisms
 - Metadata to track state for cached data
 - Controller that snoops bus (or interconnect) activity and reacts if needed to adjust the state of the cache data
- There needs to be a serialization point
 - Bus, shared L2/L3, memory controller, ...

Cache Coherence

- Informally, with coherent caches: accesses to a memory location *appear* to occur simultaneously in all copies of the memory location
 “copies” \Rightarrow caches
- Cache coherence suggests an absolute time scale -- this is not necessary
 - What is required is the "appearance" of coherence... not absolute coherence
 - E.g. temporary incoherence between memory and a write-back cache may be OK.

Write Update vs. Write Invalidate

- Coherent Shared Memory
 - All processors see the effects of others' writes
- How/when writes are propagated
 - Determine by coherence protocol

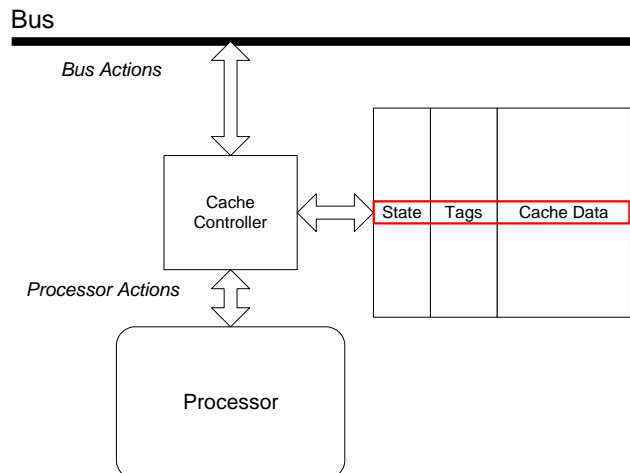


Bus-Based Snoopy Cache Coherence

- All requests broadcast on bus
- All processors and memory snoop and respond
- Cache blocks writeable at one processor or read-only at several
 - Single-writer protocol
- Snoops that hit dirty lines?
 - Flush modified data out of cache
 - Either write back to memory, then satisfy remote miss from memory, or
 - Provide dirty data directly to requestor
 - Big problem in MP systems
 - Dirty/coherence/sharing misses


Bus-Based Protocols

- Protocol consists of states and actions (state transitions)
- Actions can be invoked from processor or bus to the cache controller
- Coherence based on per cache line (block)



A Simple Protocol for Write-Through Caches

- TO SIMPLIFY ASSUME NO ALLOCATE ON STORE MISSES
 - ALL STOREs AND LOAD MISSES PROPAGATE ON THE BUS
- STOREs MAY UPDATE OR INVALIDATE CACHES
- STATE DIAGRAM
 - EACH CACHE IS REPRESENTED BY A FINITE STATE MACHINE
 - IMAGINE P IDENTICAL FSMs WORKING TOGETHER, ONE PER CACHE
 - FSM REPRESENTS THE BEHAVIOR OF A CACHE W.R.T. A MEMORY BLOCK (CACHE LINE)



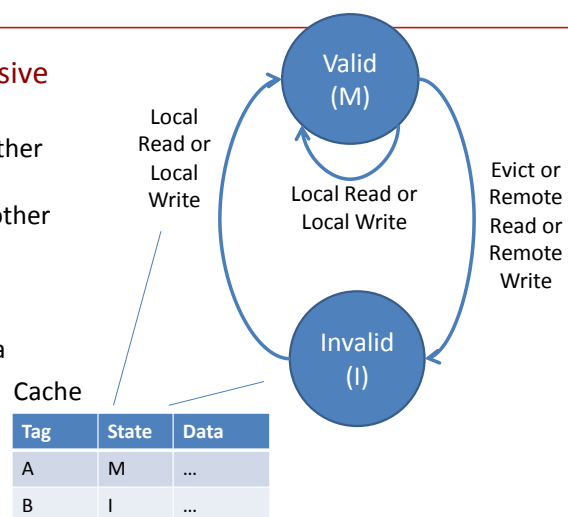
10/23/2014 (© J.P. Shen)

18-640 Lecture 14

Carnegie Mellon University 11

Minimal Coherence Protocol for Write-Back Caches

- Blocks are always private or exclusive
- State transitions:
 - Local read: I->M, fetch, invalidate other copies
 - Local write: I->M, fetch, invalidate other copies
 - Evict: M->I, write back data
 - Remote read: M->I, write back data
 - Remote write: M->I, write back data



10/23/2014 (© J.P. Shen)

18-640 Lecture 14

Carnegie Mellon University 12

Invalidate Protocol Optimization

- **Observation:** data often read shared by multiple CPUs
 - Add S (shared) state to protocol: MSI
- **State transitions:**
 - Local read: I->S, fetch shared
 - Local write: I->M, fetch modified; S->M, invalidate other copies
 - Remote read: M->I, write back data
 - Remote write: M->I, write back data

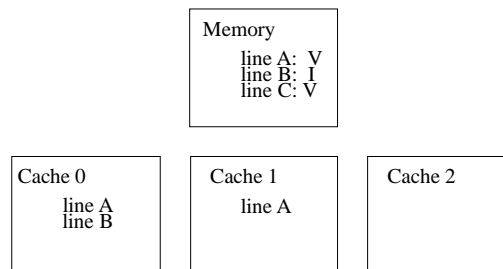
Global Coherence States

- A memory line can be present (valid) in any of the caches and/or memory
- Represent global state with an N+1 element vector
 - First N components => cache states (valid/invalid)
 - N+1st component => memory state (valid/invalid)
- **Example:**

Line A: <1,1,0,1>

Line B: <1,0,0,0>

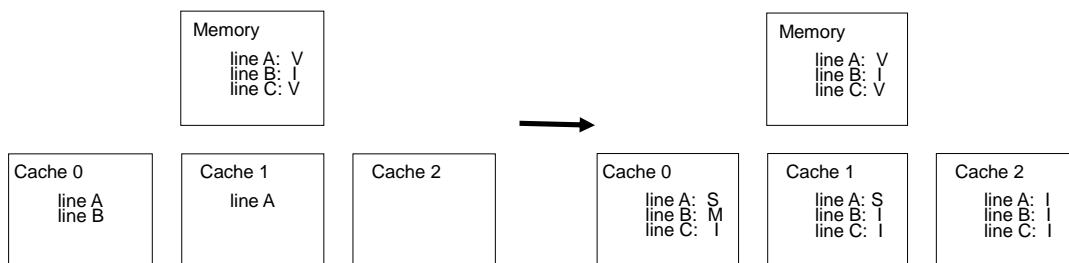
Line C: <0,0,0,1>



Local Coherence States

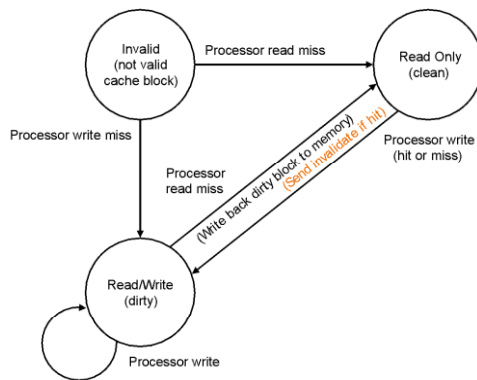
- Individual caches can maintain a summary of the state of memory lines, from a “local” perspective
 - Reduces storage for maintaining state
 - May have only partial information
- Invalid (I): $\langle 0, X, X, X, \dots, X \rangle$ -- local cache does not have a valid copy; (cache miss)
 - Don't confuse invalid state with empty frame
- Shared (S): $\langle 1, X, X, X, \dots, 1 \rangle$ -- local cache has a valid copy, main memory has a valid copy, other caches ??
- Modified (M): $\langle 1, 0, 0, \dots, 0, \dots, 0 \rangle$ -- local cache has only valid copy.
- Exclusive (E): $\langle 1, 0, 0, \dots, 0, \dots, 1 \rangle$ -- local cache has a valid copy, no other caches do, main memory has a valid copy.
- Owned (O): $\langle 1, X, X, X, \dots, X \rangle$ -- local cache has a valid copy, all other caches and memory may have a valid copy.
 - Only one cache can be in O state
 - $\langle 1, X, 1, X, \dots, 0 \rangle$ is included in O, but not included in any of the others.

Example

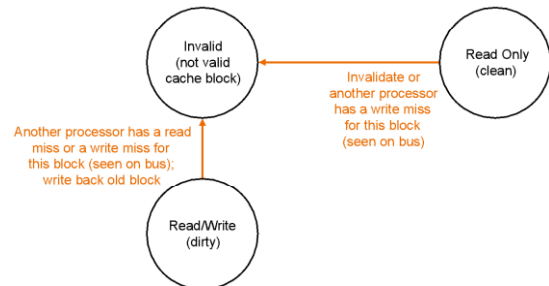


Simple Coherence Protocol FSM

[Source: Patterson/Hennessy, Comp. Org. & Design]



a. Cache state transitions using signals from the processor



b. Cache state transitions using signals from the bus

MSI Protocol

Current State	Action and Next State					
	Processor Read	Processor Write	Eviction	Cache Read	Cache Read&M	Cache Upgrade
I	Cache Read Acquire Copy → S	Cache Read&M Acquire Copy → M		No Action → I	No Action → I	No Action → I
S	No Action → S	Cache Upgrade → M	No Action → I	No Action → S	Invalidate Frame → I	Invalidate Frame → I
M	No Action → M	No Action → M	Cache Write back → I	Memory inhibit; Supply data; → S	Invalidate Frame; Memory inhibit; Supply data; → I	

MSI Example

Thread Event	Bus Action	Data From	Global State	Local States:		
				C0	C1	C2
0. Initially:			<0,0,0,1>	I	I	I
1. T0 read→	CR	Memory	<1,0,0,1>	S	I	I
2. T0 write→	CU		<1,0,0,0>	M	I	I
3. T2 read→	CR	C0	<1,0,1,1>	S	I	S
4. T1 write→	CRM	Memory	<0,1,0,0>	I	M	I

■ If line is in no cache

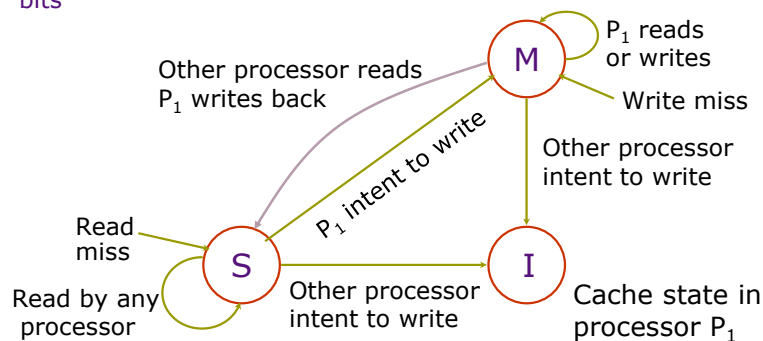
- Read, modify, Write requires 2 bus transactions
- Optimization: add Exclusive state

MSI: A Coherence Protocol for Write Back Caches

Each cache line has a tag

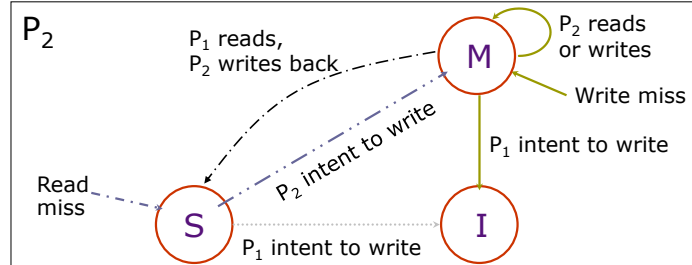
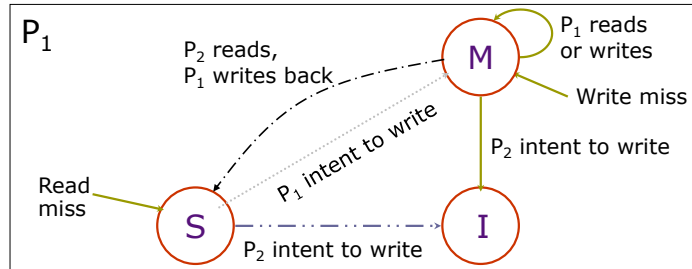


M: Modified
S: Shared
I: Invalid



MSI Coherence Protocol Example with 2 Cores

P_1 reads
 P_1 writes
 P_2 reads
 P_2 writes
 P_1 reads
 P_1 writes
 P_2 writes
 P_1 writes



10/23/2014 (© J.P. Shen)

18-640 Lecture 14

Carnegie Mellon University 21

Invalidate Protocol Optimizations

- **Observation: data can be write-private (e.g. stack frame)**
 - Avoid invalidate messages in that case
 - Add E (exclusive) state to protocol: MESI
- **State transitions:**
 - Local read: I→E if only copy, I→S if other copies exist
 - Local write: E→M silently, S→M, invalidate other copies

10/23/2014 (© J.P. Shen)

18-640 Lecture 14

Carnegie Mellon University 22

MESI Protocol

- Variation used in many Intel processors
- 4-State Protocol
 - **Modified:** $\langle 1, 0, 0, \dots, 0 \rangle$
 - **Exclusive:** $\langle 1, 0, 0, \dots, 1 \rangle$
 - **Shared:** $\langle 1, X, X, \dots, 1 \rangle$
 - **Invalid:** $\langle 0, X, X, \dots, X \rangle$
- Bus/Processor Actions
 - Same as MSI
- Adds *shared* signal to indicate if other caches have a copy

MESI Protocol

Current State	Action and Next State					
	Processor Read	Processor Write	Eviction	Cache Read	Cache Read&M	Cache Upgrade
I	Cache Read If no sharers: → E If sharers: → S	Cache Read&M → M		No Action → I	No Action → I	No Action → I
S	No Action → S	Cache Upgrade → M	No Action → I	Respond Shared: → S	No Action → I	No Action → I
E	No Action → E	No Action → M	No Action → I	Respond Shared; → S	No Action → I	
M	No Action → M	No Action → M	Cache Write-back → I	Respond dirty; Write back data; → S	Respond dirty; Write back data; → I	

MESI Example

Thread Event	Bus Action	Data From	Global State	Local States:		
				C0	C1	C2
0. Initially:			<0,0,0,1>	I	I	I
1. T0 read→	CR	Memory	<1,0,0,1>	E	I	I
2. T0 write→	none		<1,0,0,0>	M	I	I

Cache-to-Cache Transfers

- Common in many workloads:
 - T0 writes to a block: <1,0,...,0> (block in M state in T0)
 - T1 reads from block: T0 must write back, then T1 reads from memory
- In shared-bus system
 - T1 can *snarf* data from the bus during the writeback
 - Called *cache-to-cache transfer* or *dirty miss* or *intervention*
- Without shared bus
 - Must explicitly send data to requestor and to memory (for writeback)
- Known as the 4th C (cold, capacity, conflict, communication)

MESI Example 2

Thread Event	Bus Action	Data From	Global State	Local States:		
				C0	C1	C2
0. Initially:			<0,0,0,1>	I	I	I
1. T0 read→	CR	Memory	<1,0,0,1>	E	I	I
2. T0 write→	none		<1,0,0,0>	M	I	I
3. T1 read→	CR	C0	<1,1,0,1>	S	S	I
4. T2 read→	CR	Memory	<1,1,1,1>	S	S	S

MOESI Optimization

- **Observation: shared ownership prevents cache-to-cache transfer, causes unnecessary memory read**
 - Add O (owner) state to protocol: MOSI/MOESI
 - Last requestor becomes the owner
 - Avoid writeback (to memory) of dirty data
 - Also called *shared-dirty* state, since memory is stale

MOESI Protocol

- Used in AMD Opteron
- 5-State Protocol
 - Modified: <1,0,0...0>
 - Exclusive: <1,0,0,...,1>
 - Shared: <1,X,X,...,1>
 - Invalid: <0,X,X,...X>
 - Owned: <1,X,X,X,0> ; only one owner, memory not up to date
- Owner can supply data, so memory does not have to
 - Avoids lengthy memory access

MOESI Protocol

Current State	Action and Next State					
	Processor Read	Processor Write	Eviction	Cache Read	Cache Read&M	Cache Upgrade
<i>I</i>	Cache Read If no sharers: → E If sharers: → S	Cache Read&M → M		No Action → I	No Action → I	No Action → I
<i>S</i>	No Action → S	Cache Upgrade → M	No Action → I	Respond shared; → S	No Action → I	No Action → I
<i>E</i>	No Action → E	No Action → M	No Action → I	Respond shared; Supply data; → S	Respond shared; Supply data; → I	
<i>O</i>	No Action → O	Cache Upgrade → M	Cache Write- back → I	Respond shared; Supply data; → O	Respond shared; Supply data; → I	
<i>M</i>	No Action → M	No Action → M	Cache Write- back → I	Respond shared; Supply data; → O	Respond shared; Supply data; → I	

MOESI Example

Thread Event	Bus Action	Data From	Global State	local states		
				C0	C1	C2
0. Initially:			<0,0,0,1>	I	I	I
1. T0 read→	CR	Memory	<1,0,0,1>	E	I	I
2. T0 write→	none		<1,0,0,0>	M	I	I
3. T2 read→	CR	C0	<1,0,1,0>	O	I	S
4. T1 write→	CRM	C0	<0,1,0,0>	I	M	I

MOESI Coherence Protocol

- A protocol that tracks validity, ownership, and exclusiveness
 - Modified: dirty and private
 - Owned: dirty but shared
 - Avoid writeback to memory on M->S transitions
 - Exclusive: clean but private
 - Avoid upgrade misses on private data
 - Shared
 - Invalid
- There are also some variations (MOSI and MESI)
- What happens when 2 cores read/write different words in a cache line?

Snooping with Multi-level Caches

■ Private L2 caches

- If inclusive, snooping traffic checked at the L2 level first
- Only accesses that refer to data cached in L1 need to be forwarded
- Saves bandwidth at the L1 cache

■ Shared L2 or L3 caches

- Can act as serialization points even if there is no bus
- Track state of cache line and list of sharers (bit mask)
- Essentially the shared cache acts like a coherence directory

Scaling Coherence Protocols

■ The problem

- Too much broadcast traffic for snooping (probing)

■ Solution: probe filters

- Maintain info of which address ranges that are definitely not shared or definitely shared
- Allows filtering of snoop traffic

■ Solution: directory based coherence

- A directory stores all coherence info (e.g., sharers)
- Consult directory before sending coherence messages
- Caching/filtering schemes to avoid latency of 3-hops

Implementing Cache Coherence

- **Snooping implementation**
 - Origins in shared-memory-bus systems
 - All CPUs could observe all other CPUs requests on the bus; hence “snooping”
 - Bus Read, Bus Write, Bus Upgrade
 - React appropriately to snooped commands
 - Invalidate shared copies
 - Provide up-to-date copies of dirty lines
 - Flush (writeback) to memory, or
 - Direct intervention (*modified intervention* or *dirty miss*)
- **Snooping suffers from:**
 - Scalability: shared busses not practical
 - Ordering of requests without a shared bus
 - Lots of recent and on-going work on scaling snoop-based systems

Snooping Cache Coherence

- **Basic idea: broadcast snoop to all caches to find owner**
- **Not scalable?**
 - Address traffic roughly proportional to square of number of processors
 - Current implementations scale to 64/128-way (Sun/IBM) with multiple address-interleaved broadcast networks
- **Inbound snoop bandwidth: big problem**

$$OutboundSnoopRate = s_o = \langle CacheMissRate \rangle + \langle BusUpgradeRate \rangle$$

$$InboundSnoopRate = s_i = n \times s_o$$

Snoop Bandwidth

- Snoop filtering of various kinds is possible
- Filter snoops at sink: Jetty filter [Moshovos et al., HPCA 2001]
 - Check small “filter cache” that summarizes contents of local cache
 - Avoid power-hungry lookups in each tag array
- Filter snoops at source: Multicast snooping [Bilir et al., ISCA 1999]
 - Predict likely sharing set, snoop only predicted sharers
 - Double-check at directory to make sure
- Filter snoops at source: Region coherence
 - Concurrent work: [Cantin/Smith/Lipasti, ISCA 2005; Moshovos, ISCA 2005]
 - Check larger region of memory on every snoop; remember when no sharers
 - Snoop only on first reference to region, or when region is shared
 - Eliminate 60%+ of all snoops

Snoop Latency

- Snoop latency:
 - Must reach all nodes, return and combine responses
 - Topology matters: ring, mesh, torus, hypercube
 - No obvious solutions
- Parallelism: fundamental advantage of snooping
 - Broadcast exposes parallelism, enables speculative latency reduction

Scaleable Cache Coherence

- No physical bus but still snoop
 - Point-to-point tree structure (indirect) or ring
 - Root of tree or ring provide ordering point
 - Use some scalable network for data (ordering less important)
- Or, use level of indirection through directory
 - Directory at memory remembers:
 - Which processor is “single writer”
 - Which processors are “shared readers”
 - Level of indirection has a price
 - Dirty misses require 3 hops instead of two
 - Snoop: Requestor->Owner->Requestor
 - Directory: Requestor->Directory->Owner->Requestor

Implementing Cache Coherence

- Directory implementation
 - Extra bits stored in memory (directory) record state of line
 - Memory controller maintains coherence based on the current state
 - Other CPUs' commands are not snooped, instead:
 - Directory forwards relevant commands
 - Powerful filtering effect: only observe commands that you need to observe
 - Meanwhile, bandwidth at directory scales by adding memory controllers as you increase size of the system
 - Leads to very scalable designs (100s to 1000s of CPUs)
- Directory shortcomings
 - Indirection through directory has latency penalty
 - If shared line is dirty in other CPU's cache, directory must forward request, adding latency
 - This can severely impact performance of applications with heavy sharing (e.g. relational databases)

Directory Protocol Implementation


- Basic idea: Centralized directory keeps track of data location(s)
- Scalable
 - Address traffic roughly proportional to number of processors
 - Directory & traffic can be distributed with memory banks (interleaved)
 - Directory cost (SRAM) or latency (DRAM) can be prohibitive
- Presence bits track sharers
 - Full map (N processors, N bits): cost/scalability
 - Limited map (limits number of sharers)
 - Coarse map (identifies board/node/cluster; must use broadcast)
- Vectors track sharers
 - Point to shared copies
 - Fixed number, linked lists (SCI), caches chained together
 - Latency vs. cost vs. scalability

Directory Protocol Latency

- Access to non-shared data
 - Overlap directory read with data read
 - Best possible latency given distributed memory
- Access to shared data
 - Dirty miss, modified intervention
 - Shared intervention?
 - If DRAM directory, no gain
 - If directory cache, possible gain (use F state)
 - No inherent parallelism
 - Indirection adds latency
 - Minimum 3 hops, often 4 hops

Directory-Based Cache Coherence

- An alternative for large, scalable MPs
- Can be based on any of the protocols discussed thus far
 - We will use MSI
- Memory Controller becomes an active participant
- Sharing info held in memory directory
 - Directory may be distributed
- Use point-to-point messages
- Network is not totally ordered

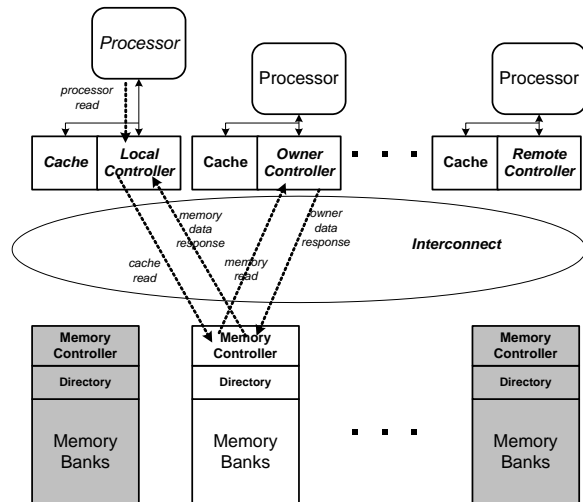


Example: Simple Directory Protocol

- Local cache controller states
 - M, S, I as before
- Local directory states
 - **Shared**: $\langle 1, X, X, \dots, 1 \rangle$; one or more proc. has copy; memory is up-to-date
 - **Modified**: $\langle 0, 1, 0, \dots, 0 \rangle$ one processor has copy; memory does not have a valid copy
 - **Uncached**: $\langle 0, 0, \dots, 0, 1 \rangle$ none of the processors has a valid copy
- Directory also keeps track of sharers
 - Can keep global state vector in full
 - e.g. via a bit vector

Example

- *Local cache* suffers load miss
- Line in *remote cache* in M state
 - It is the *owner*
- Four messages send over network
 - Cache read from local controller to home memory controller
 - Memory read to remote cache controller
 - Owner data back to memory controller; change state to S
 - Memory data back to local cache; change state to S



Cache Controller State Table

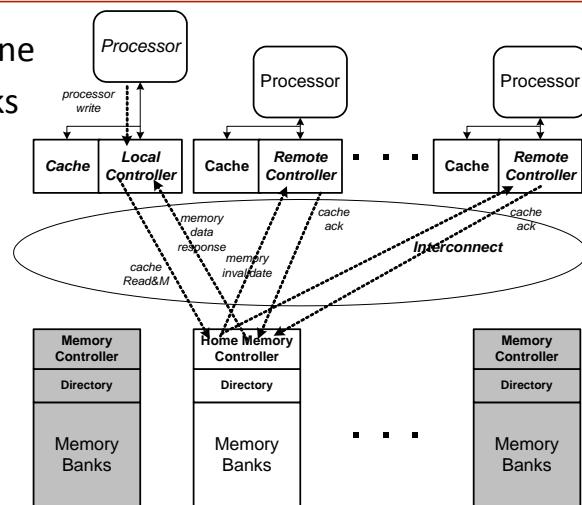
	Cache Controller Actions and Next States							
	from Processor Side			from Memory Side				
Current State	Processor Read	Processor Write	Eviction	Memory Read	Memory Read&M	Memory Invalidate	Memory Upgrade	Memory Data
I	Cache Read → I'	Cache Read&M → I''				No Action → I		
S	No Action → S	Cache Upgrade → S'	No Action* → I			Invalidate Frame; Cache ACK; → I		
M	No Action → M	No Action → M	Cache Write-back → I	Owner Data; → S	Owner Data; → I	Invalidate Frame; Cache ACK; → I		
I'								Fill Cache → S
I''								Fill Cache → M
S'							No Action → M	

Memory Controller State Table

	Memory Controller Actions and Next States						
	command from Local Cache Controller				response from Remote Cache Controller		
Current Directory State	Cache Read	Cache Read&M	Cache Upgrade		Data Write-back	Cache ACK	Owner Data
U	Memory Data; Add Requestor to Sharers; → S	Memory Data; Add Requestor to Sharers; → M					
S	Memory Data; Add Requestor to Sharers; → S	Memory Invalidate All Sharers; → M'	Memory Upgrade All Sharers; → M''		No Action → I		
M	Memory Read from Owner; → S'	Memory Read&M; to Owner → M'			Make Sharers Empty; → U		
S'							Memory Data to Requestor; Write memory; Add Requestor to Sharers; → S
M'						When all ACKS Memory Data; → M	Memory Data to Requestor; → M
M''						When all ACKS then → M	

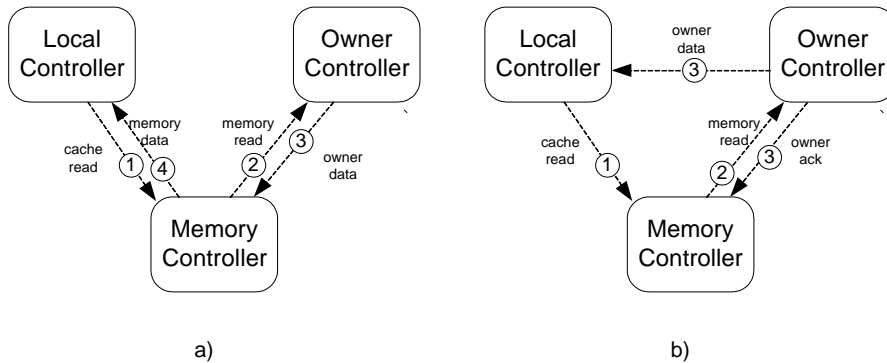
Another Example

- Local write (miss) to shared line
- Requires invalidations and acks



Variation: Three Hop Protocol

- Have owner send data directly to local controller
- Owner Acks to Memory Controller in parallel



10/23/2014 (© J.P. Shen)

18-640 Lecture 14

Carnegie Mellon University 49

Example Sequence

- Similar to earlier sequences

Thread Event	Controller Actions	Data From	global state	local states: C0 C1 C2		
0. Initially:			<0,0,0,1>	I	I	I
1. T0 read→	CR,MD	Memory	<1,0,0,1>	S	I	I
2. T0 write→	CU, MU*,MD		<1,0,0,0>	M	I	I
3. T2 read→	CR,MR,MD	C0	<1,0,1,1>	S	I	S
4. T1 write→	CRM,MI,CA,MD	Memory	<0,1,0,0>	I	M	I

10/23/2014 (© J.P. Shen)

18-640 Lecture 14

Carnegie Mellon University 50

Directory Protocol Optimizations

- Remove dead blocks from cache:
 - Eliminate 3- or 4-hop latency
 - Dynamic Self-Invalidation [Lebeck/Wood, ISCA 1995]
 - Last touch prediction [Lai/Falsafi, ISCA 2000]
 - Dead block prediction [Lai/Fide/Falsafi, ISCA 2001]
- Predict sharers
 - Prediction in coherence protocols [Mukherjee/Hill, ISCA 1998]
 - Instruction-based prediction [Kaxiras/Goodman, ISCA 1999]
 - Sharing prediction [Lai/Falsafi, ISCA 1999]
- Hybrid snooping/directory protocols
 - Improve latency by snooping, conserve bandwidth with directory
 - Multicast snooping [Bilir et al., ISCA 1999; Martin et al., ISCA 2003]
 - Bandwidth-adaptive hybrid [Martin et al., HPCA 2002]
 - Token Coherence [Martin et al., ISCA 2003]
 - Virtual Tree Coherence [Enright Jerger MICRO 2008]

Update Protocols

- **Basic idea:**
 - All writes (updates) are made visible to all caches:
 - (address,value) tuples sent “everywhere”
 - Similar to write-through protocol for uniprocessor caches
 - Obviously not scalable beyond a few processors
 - No one actually builds machines this way
- **Simple optimization**
 - Send updates to memory/directory
 - Directory propagates updates to all known copies: less bandwidth
- **Further optimizations: combine & delay**
 - Write-combining of adjacent updates (if consistency model allows)
 - Send write-combined data
 - Delay sending write-combined data until requested
- **Logical end result**
 - Writes are combined into larger units, updates are delayed until needed
 - Effectively the same as invalidate protocol
- **Of historical interest only (Firefly and Dragon protocols)**

Update vs Invalidate

- [Weber & Gupta, ASPLOS3]
 - Consider sharing patterns
- No Sharing
 - Independent threads
 - Coherence due to thread migration
 - Update protocol performs many wasteful updates
- Read-Only
 - No significant coherence issues; most protocols work well
- Migratory Objects
 - Manipulated by one processor at a time
 - Often protected by a lock
 - Usually a write causes only a single invalidation
 - E state useful for Read-modify-Write patterns
 - Update protocol could proliferate copies

Update vs Invalidate, contd.

- Synchronization Objects
 - Locks
 - Update could reduce spin traffic invalidations
 - Test & Test&Set w/ invalidate protocol would work well
- Many Readers, One Writer
 - Update protocol may work well, but writes are relatively rare
- Many Writers/Readers
 - Invalidate probably works better
 - Update will proliferate copies
- What is used today?
 - Invalidate is dominant
 - CMP *may* change this assessment
 - more on-chip bandwidth