

18-640 Foundations of Computer Architecture

Lecture 16: “SIMD and GPU Processors”

Mridula Chappalli Srinivasa
(Ken Lueh, PhD - Sr. Principal Engineer, Intel)
October 30, 2014

➤ Required Reading Assignments:

- Chapter 4 from new Hennessy & Patterson textbook
 - Computer Architecture, 5th edition
 - Read 4.4 – 4.10 for today's lecture

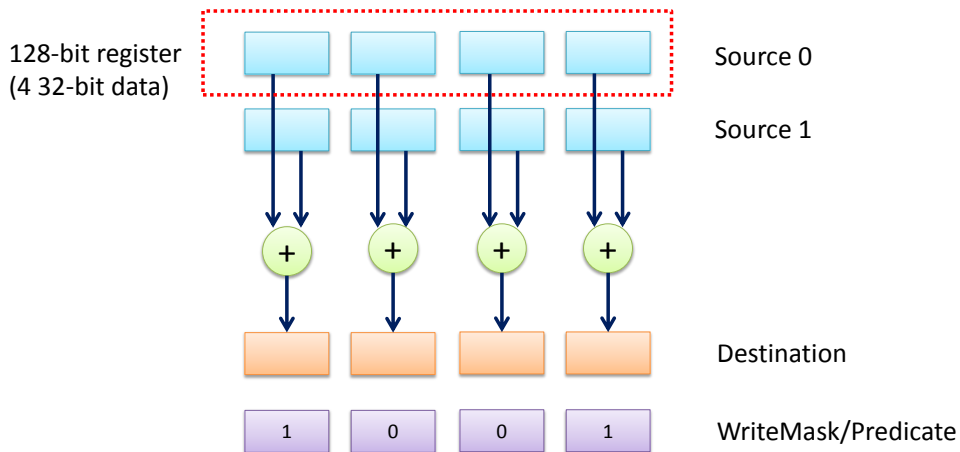


Carnegie Mellon University ¹

10/30/2014 (© J.P. Shen)

Lecture 16

Single Instruction Multiple Data (SIMD)



10/30/2014 (© J.P. Shen)

Lecture 16

Carnegie Mellon University ²

SIMD Extensions for Superscalar Processors


- Every CISC/RISC processor today has SIMD extensions
 - MMX, SSE, SSE-2, SSE-3, SSE-4, AVX, AVX2, AltiVec, VIS, ...
- Basic idea: accelerate multimedia processing
 - Define vectors of 8, 16, 32 and 64 bit elements in regular registers
 - Apply SIMD arithmetic on these vectors
- Nice and cheap
 - Don't need to define big vector register file
 - This has changed in more recent SIMD extensions
 - All we need to do
 - Add the proper opcodes for SIMD arithmetic
 - Modify datapaths to execute SIMD arithmetic
 - Certain operations are easier on short vectors
 - Reductions, random permutations

Problems with SIMD Extension

- SIMD defines short, fixed-sized, vectors
 - Cannot capture data parallelism wider than 64 bits
 - MMX (1996) has 64-bit registers (8 8-bit or 4 16-bit operations)
 - Must use wide-issue to utilize more than 64-bit datapaths
 - SSE and AltiVec have switched to 128-bits because of this
 - AVX2 has switched to 512-bits because of this
- SIMD does not support vector memory accesses
 - Strided and indexed accesses for narrow elements
 - Needs multi-instruction sequence to emulate
 - Pack, unpack, shift, rotate, merge, etc
 - Cancels most of performance and code density benefits of vectors
- Compiler support for SIMD?
 - Auto vectorization is hard
 - Reply on programming model (e.g., OpenMP, Cilk+)


AVX2 SIMD Register Set

- Intel® AVX extends all 16 XMM registers to 256bits
- Intel AVX instructions operate on either:
 - The whole 256-bits (FP only)
 - The lower 128-bits (like existing Intel® SSE instructions)
 - A replacement for existing scalar/128-bit SSE instructions
 - Provides new capabilities on existing instructions
 - The upper 128-bits of the register are zeroed out
- Intel AVX2 supports integer operations




AVX Example

- Intel® AVX defines two 128-bit lanes (low =xmm, high=ymm[255:128])
 - Nearly all operations are defined as "in-lane"
 - For most instructions, e.g., VADDPS, the lane division is uninteresting



- Some in-lane behavior is more interesting: VPERMILPS



Intel® SSE functionality is preserved within lanes

Graphics Processors Timeline

- Till mid 90s
 - VGA controllers used to accelerate some display functions
- Mid 90s to mid 00s
 - Fixed-function graphics accelerators for the OpenGL and DirectX APIs
 - Some GP-GPU capabilities by on top of the interfaces
 - 3D graphics: triangle setup & rasterization, texture mapping & shading
- Modern GPUs
 - Programmable multiprocessors (optimized for data-parallel ops)
 - OpenGL, DirectX, C++ AMP, OpenCL, RenderScript, Cilk+
 - Some fixed function hardware (texture, raster ops, ...)

Why GPU?

Theoretical GFLOP/s





Fig 2: NVIDIA CUDA C Programming Guide

CPU vs GPU

Fig 3: NVIDIA
CUDA
Programming
Guide

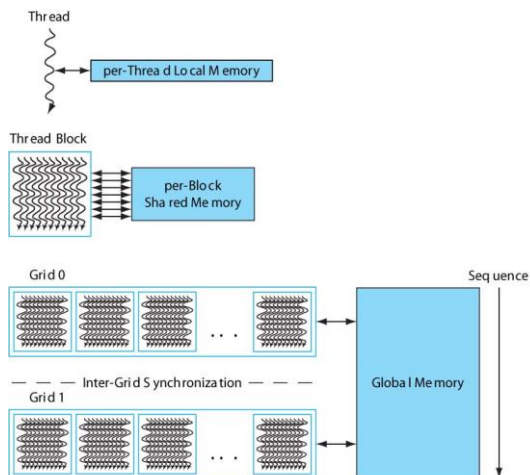


10/30/2014 (© J.P. Shen)

Lecture 16

Carnegie Mellon University 9

GPU Thread Model - Software View



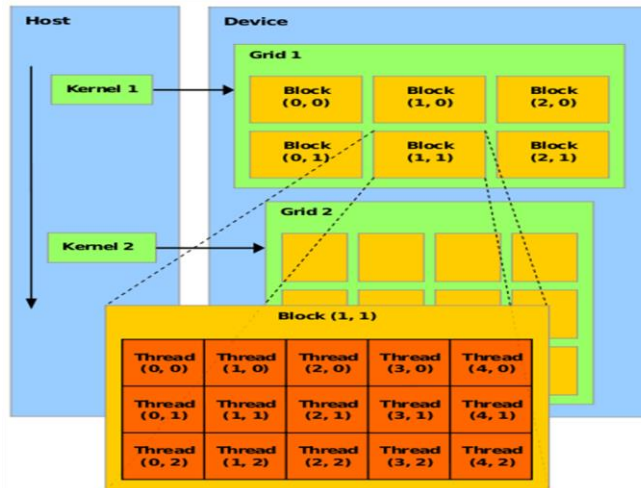
- Single-program multiple data (SPMD) model
- Each thread has local memory
- Parallel threads packed in blocks
 - Access to per-block shared memory
 - Can synchronize with barrier
- Grids include independent groups
 - May execute concurrently

10/30/2014 (© J.P. Shen)

Lecture 16

Carnegie Mellon University 10

Data Parallel Decomposition



- Programmer partitions problems into grids and blocks
- Programmer partitions blocks into threads
- Built-in variables
 - blockIdx.x, blockIdx.y, blockIdx.z
 - threadIdx.x, threadIdx.y, threadIdx.z
 - blockDim.x, blockDim.y, blockDim.z
 - gridDim.x, gridDim.y, gridDim.z

10/30/2014 (© J.P. Shen)

Lecture 16

Carnegie Mellon University 11

Code Example: DAXPY

C Code

```
// Invoke DAXPY
daxpy(n, 2.0, x, y);
// DAXPY in C
void daxpy(int n, double a, double *x, double *y)
{
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}
```

CUDA Code

```
// Invoke DAXPY with 256 threads per block
__host__
int nblocks = (n+ 255) / 256;
daxpy<<<nblocks, 256>>>(n, 2.0, x, y);
// DAXPY in CUDA
__device__
void daxpy(int n, double a, double *x, double *y)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < n) y[i] = a*x[i] + y[i];
}
```

- CUDA code launches 256 threads per block
 - Thread = 1 iteration of scalar loop (1 element in vector loop)
 - Block = body of vectorized loop (with VL=256 in this example)
 - Grid = vectorizable loop

10/30/2014 (© J.P. Shen)

Lecture 16

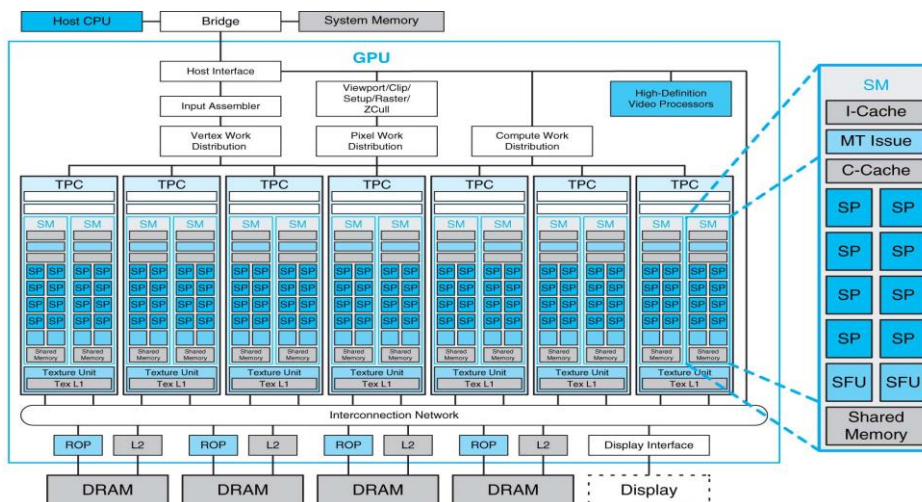
Carnegie Mellon University 12

Sample CUDA Program

```
__global__ void vcos( int n, float* x, float* y ) {
    int ix = blockIdx.x*blockDim.x + threadIdx.x; }

int main() {
    float *host_x, *host_y;
    float *dev_x, *dev_y;
    int n = 1024;
    host_x = (float*)malloc( n*sizeof(float) );
    host_y = (float*)malloc( n*sizeof(float) );
    cudaMalloc( &dev_x, n*sizeof(float) );
    cudaMalloc( &dev_y, n*sizeof(float) );
    /*TODO - Fill x[i] with data */
    cudaMemcpy( dev_x, host_x, n*sizeof(float), cudaMemcpyHostToDevice );
    bk = (int)( n / 256 );
    vcos<<<bk,256>>>( n, dev_x, dev_y );
    cudaMemcpy( host_y, dev_y, n*sizeof(float), cudaMemcpyDeviceToHost );
    return(0) }
```

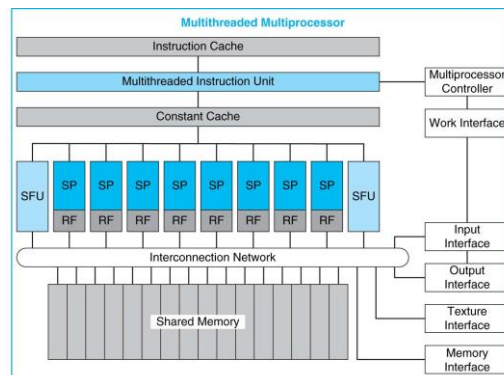
GPU Architecture: Nvidia GeForce 8800 (aka Tesla Architecture)



GPU Architecture

- A highly multithreaded, multiprocessor system
 - 100s of streaming processors (SPs)
 - 8 SPs in a streaming multiprocessor (SM) with some caches
 - 2 SMs in a texture processor cluster (TPCs) with one texture pipe
- Scheduling controlled mostly by hardware
- Scalability
 - By scaling the number of TPCs and memory channels
- Fixed function components for graphics
 - Texture pipes and caches, raster operation units (ROP), ...

Streaming Multiprocessor

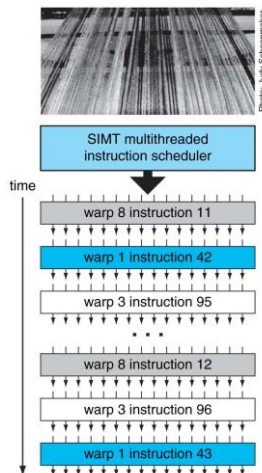


- Another way to think of this
 - SM = a multithreaded vector core
 - SP (streaming processor) = a vector or SIMD lane (registers + ALUs)

Streaming Multiprocessor Details

- Each SP is a simple processor core
 - 1024 32-bit registers shared flexibly by up to 64 threads
 - Integer and floating-point arithmetic units
 - Including multiply add
- Special function unit
 - Implements functions such as divide, square root, sine, cosine, ...
- Instruction cache and constants cache
 - Shared by all threads
- Multibanked shared memory
 - E.g. 16 banks to allow parallel accesses by all SP

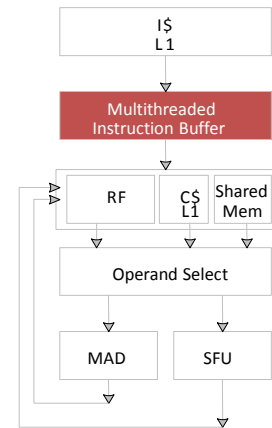
Instruction & Thread Scheduling: Where Threads Meets Data Parallelism



- In theory, all threads can be independent
 - SM hardware implements zero-overhead switching
 - 768 threads (24 warps)
- For efficiency, 32 threads are packed in warps
 - Warp: set of parallel threads that execute same instruction
 - Warp = a thread of vector instructions
 - Warps introduce data parallelism
 - 1 warp instruction keeps SPs busy for 4 cycles in Tesla
- Individual threads may be inactive
 - Because they branched differently or predication
 - This is the equivalent of conditional execution
 - Loss of efficiency if not data parallel
- Software thread blocks mapped to warps
 - When HW resources are available

Instruction Buffering & Warp Scheduling

- Fetch one instruction/cycle
 - From the L1 instruction cache into an instruction buffer slot
- Issue one “ready-to-go” instruction/cycle
 - All elements of the warp must be ready
 - Scoreboarding used to track hazards and determine ready warps
 - Round-robin or age based selection between ready warps
- Instruction broadcasted to all SP
 - Will keep SPs busy for up to 4 cycles



Dependence Tracking Using a Scoreboard

- Status of all register operands is tracked
 - RAW hazards for high-latency operations
 - Dependencies to memory operations
- Instructions are ready when all register operands are ready for whole warp
 - Divergence of threads within warps is also tracked
- A warp may be blocked because of
 - Dependencies through registers
 - Synchronization operations (barriers or atomic ops)
 - Memory accesses
- But other warps can proceed in order to hide latency

Tracking Branch Divergence

- Similar to vector processors but masks handled internally
 - No explicit mask register
 - Per warp stack that stores PCs and masks for “not taken” paths
- On a conditional branch
 - Push the current mask onto stack
 - Push the mask and PC for the “not taken” path
 - Set the mask for the “taken” path and execute
- At the end of the “taken” path
 - Pop the mask and PC for the “not taken” path and execute
- At the end of the “not taken” path
 - Pop the original mask before the branch instruction
- If a mask is all zeros, the block is skipped

Tracking Branch Divergence

Mask = 1111	→	A = Data[i]; B = Data[j];
Push Mask = 1100 Set Mask = 0011	→	If (condition A) {
Push Mask = 0010 Set Mask = 0001	→	if (condition B) {
Pop Mask = 0010	→	Y = A+B; } else { Y = A*B; }
Pop Mask = 1100	→	} else { Y = A-B; }
Mask = 1111	→	}


Memory System

- Per SM caches/memories
 - Instruction and constant caches
 - Multi-banked shared memory
- Distributed texture cache
 - Per TPC L1 and distributed L2 cache
 - Specialized for texture accesses in graphics pipeline
 - Moving towards generalized and shared L2 in upcoming chips
- Multi-channel DRAM main memory (e.g. 8 DDR-3 channels)
 - Interleaved addresses to achieve higher bandwidth
 - Lossless and lossy compression used to increase bandwidth
 - Aggressive access scheduling used to increase bandwidth
- Per thread private memory and global memory mapped to DRAM
 - Relying on threads to hide long latencies

Memory Access Coalescing

Goal: Combine multiple memory accesses generated from multiple threads into a single physical transaction


Paulius Micikeviciu,
"M02: High Performance
Computing with CUDA:
Optimizing CUDA",
SC08, Austin, TX.



Memory Bank Conflicts

A bank conflict occurs if two or more threads access any bytes within different 32-bit words belonging to the same bank.

NVIDIA CUDA C Programming
Guide Version 4.0, Figure F-2



Synchronization

- Barrier synchronization within a thread block
 - Tracking simplified by grouping threads into warps
 - Counter used to track number of threads that have arrived to barrier
- Atomic operations to global memory
 - Atomic read-modify-write (add, min, max, and, or, xor)
 - Atomic exchange or compare and swap
 - They are tied to DRAM latency
 - Moving to shared L2 in upcoming chips

Intel Graphics ISA Syntax

• **[Predicate] Instr (ExecSize) Dst Src0 Src1 {Statement}**

- Predicate: +, -, Channel Replication (+x)
- Instr: opcode[.modifier][.flagmod]
 - Instruction modifier: add.sat
 - Flag modifier: mul.f0.g; add.f1.e
- Dst: RegFile RegNum<Region>.dmask:type
- Src0: (SrcMod) RegFile RegNum<Region>.swizzle:type
 - Source Modifier (SrcMod): -, (abs), -(abs)
- Src1: (SrcMod) RegFile RegNum<Region>.swizzle:type
- Example
 - **(+f0.0) add.f0.g (8) r2<8;8,1>:f -r10<8;8,1>:f -(abs) r12<8;8,1>:f**
 - Details later...

Intel Graphics ISA

- Arithmetic
 - add, mul, mad, mac, ...
- Bitwise
 - and, or, not, xor, shl, shr, bfe, ...
- Control flow
 - if, else, endif, while, call, break, brc, jmp, ...
- Math
 - Inv, sin, cos, idiv, log, ...
- Atomic
 - And, or, add, sub, inc, cmpwr,
- Barrier
- Load/store
- Sample
- Misc


SIMD Control Flow

- HW support for SIMD control flow
- Execution paths of multiple threads start diverging
- Efficiency of throughput drops ...

```

cmp.ne.f0.0 (8) null<1>:w r26.0<16;16,1>:w 0
(+f0) If (8) off_endif off_else
...
else (8) off_endif
...
endif (8) off_redirect
  
```

Active channels  Inactive channels 



10/30/2014 (© J.P. Shen)

Lecture 16

Carnegie Mellon University 29

SIMD while

LOOP_START:

```

...
cmp.ne.f0.0 (8) null<1>:w r26.0<16;16,1>:w 0:w
(+f0.0) if (8) IF_LABEL__1 IF_LABEL__1
break (8) IF_LABEL__1 LOOP_BREAK
  
```

IF_LABEL__1:
endif (8)

```

...
cmp.g.f0.0 (8) null<1>:d r2.0<8;8,1>:d 0:w
  
```

LOOP_BREAK:

(+f0.0) while (8) LOOP_START

Active channels  Inactive channels 



Active channels IPs is set to
LOOP_BREAK

10/30/2014 (© J.P. Shen)

Lecture 16

Carnegie Mellon University 30

SIMD while

```

LOOP_START:  ←-----
    ...
    cmp.ne.f0.0 (8) null<1>:w r26.0<16;16,1>:w 0:w
    (+f0.0) if (8) IF_LABEL__1 IF_LABEL__1
        break (8) IF_LABEL__1 LOOP_BREAK
IF_LABEL__1:
    endif (8)
    ...
    cmp.g.f0.0 (8) null<1>:d r2.0<8;8,1>:d 0:w
LOOP_BREAK:
    (+f0.0) while (8) LOOP_START
  
```

Reactivate channels at endif



Active channels IPs are set to
LOOP_START



10/30/2014 (© J.P. Shen)

Lecture 16

Carnegie Mellon University 31

SIMD while

```

LOOP_START:
    ...
    cmp.ne.f0.0 (8) null<1>:w r26.0<16;16,1>:w 0:w
    (+f0.0) if (8) IF_LABEL__1 IF_LABEL__1
        break (8) IF_LABEL__1 LOOP_BREAK
    endif (8)
IF_LABEL__1:
    ...
    cmp.g.f0.0 (8) null<1>:d r2.0<8;8,1>:d 0:w
LOOP_BREAK:
    (+f0.0) while (8) LOOP_START  ←-----
  
```

Active channels IPs are set to
LOOP_START



All channels are inactive



Reactivate channels and
Advance to IP+1



10/30/2014 (© J.P. Shen)

Lecture 16

Carnegie Mellon University 32

Fundamental Challenges of GPGPU

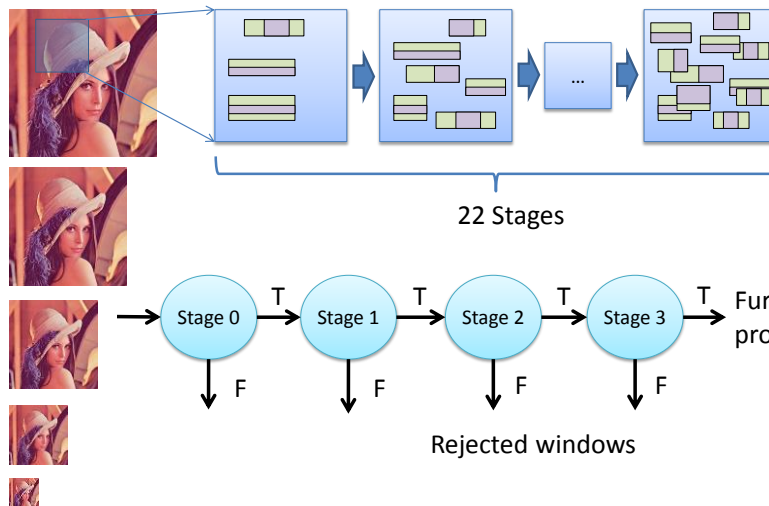
- Diverging control flow
- Irregular memory access
- Load balancing between CPU and GPU

10/30/2014 (© J.P. Shen)

Lecture 16

Carnegie Mellon University 33

Case Study: Viola-Jones Face Detector



- Based on Harr-like feature
- 24x24 detecting window
- Scaling is 1.2
- Check feature against threshold
- Step 2 pixels
 - 512x512 image needs ~60K sub windows

10/30/2014 (© J.P. Shen)

Lecture 16

Carnegie Mellon University 34

Diverging Control Flow

Pseudo Code:


```

for i in n_stages:
    for f in stage_features[i]:
        sum += weakClassifier(f)

    If sum < stage_threshold[i]:
        return False
return True

```

SIMD8: 8 sub windows



Efficiency drops to 12.5%



Stage	1	2	3	4	5	6	...	20	21	22
HAAR Feature Count	3	16	21	39	33	44		182	211	213
Eliminate Averagely (%)	20.41	30.04	25.84	11.95	3.32	3.07		0.02	0.01	0.01

First 3 stages eliminate 75+% sub windows on average



10/30/2014 (© J.P. Shen)

Lecture 16

Carnegie Mellon University 35


Divide Into Phases & Exploit Different Parallelism

Phase 1: exploit parallelism of sub windows



Gather positive sub windows

Phase 2: exploit parallelism of features



- Each stage has enough features
- Iterate through 8 features at one time
- High efficiency

10/30/2014 (© J.P. Shen)

Lecture 16

Carnegie Mellon University 36

Reference

- HSW - Volume 7: 3D Media GPGPU
 - <https://01.org/linuxgraphics/documentation/2013-intel-core-processor-family>
- NVIDIA TESLA: A UNIFIED GRAPHICS AND COMPUTING ARCHITECTURE
 - http://www.cs.cmu.edu/afs/cs/academic/class/15869-f11/www/readings/lindholm08_tesla.pdf
- NVIDIA CUDA C Programming Guide
 - <http://docs.nvidia.com/cuda/cuda-c-programming-guide#axzz3HbLR7y5t>